

Deep Hedging

Ben Wood

Quantitative Research

JPMorgan

Oxford ML x Finance Summer School

August 2022

Overview

Derivatives trading

- Derivatives
- What a trader does
- Hedging as reinforcement learning

Deep Hedging

- Rewards and utility functions
- Optimized certainty equivalents
- Architecture
- Examples

Market simulation

- Option prices
- Compression
- Time series generation

Deep Bellman Hedging

- Beyond policy search

Derivatives primer

Derivatives

- Financial contracts defining payments derived from the prices of underlying assets
- Stocks, indices, bonds, rates, FX, commodities

Markets

- Large in volume and notional terms
- Standard derivatives are traded on exchanges
- More complex derivatives are traded directly between counterparties (OTC)

Participants

- Sell side: banks, market making firms
- Buy side: asset managers, hedge funds, pension funds, insurance companies, retail investors

Activities

- Investment
- Hedging
- Market making

Examples

FX futures contract

- Contract to exchange fixed amounts of two currencies on a given future date
 - British Pound Futures Sep 22 1.214

Equity index call option

- Contract that pays the amount by which the index level at **maturity** exceeds the **strike price**, if that amount is positive: $V_T = \max(0, I_T - K)$
 - Euro Stoxx 50 Dec 22 3800 CALL

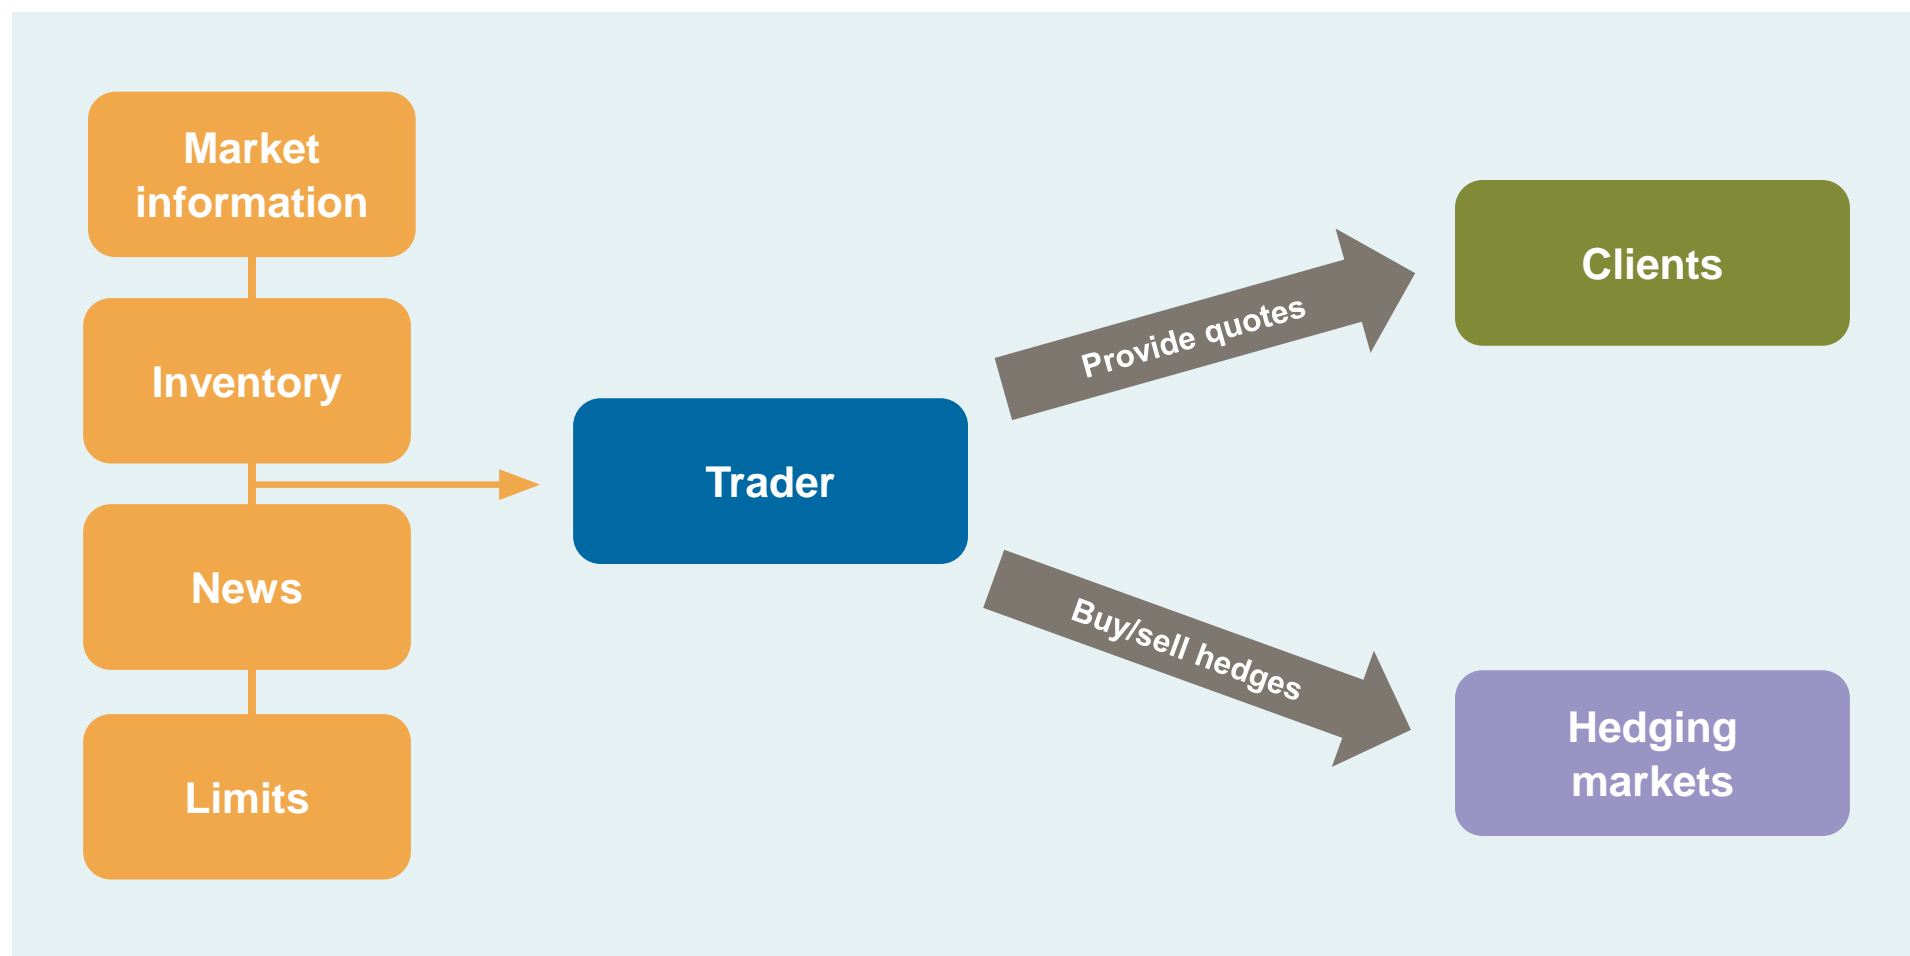
Equity worst-of basket autocallable

- Exotic option tracking the performance of the worst-performing stock in a basket
- Pays a quarterly coupon if the worst-of is above a coupon barrier threshold at quarter end
- Terminates early if the worst-of is above a higher knockout barrier threshold at quarter end
- Repays the notional on termination
- Repays the notional minus a down-and-in put at maturity

Derivatives trading

What does a (sell-side) trader do?

- Key tasks are quoting and hedging
- Monitor the market, know her inventory, react to news, understand risk and PnL, make sure to comply with controls and limits



Trading derivatives involves risk

- At least one future payment of uncertain amount
- How can we reduce the risk?
- Trade in the underlying asset

Example: forward contract

- At maturity T , we pay $V_T = S_T - K$
- To reduce the risk, we can simply borrow money and buy the stock today
- All the uncertainty has been removed

Example: call option

- At maturity T , we pay $V_T = N \max(0, S_T - K)$
- To reduce the risk, we can borrow money and buy the stock today
- We will need somewhere between 0 and N units
- If the stock price goes up, we need to buy more, and if it goes down, we need to sell



Classical models

Risk-neutral valuation models from math finance

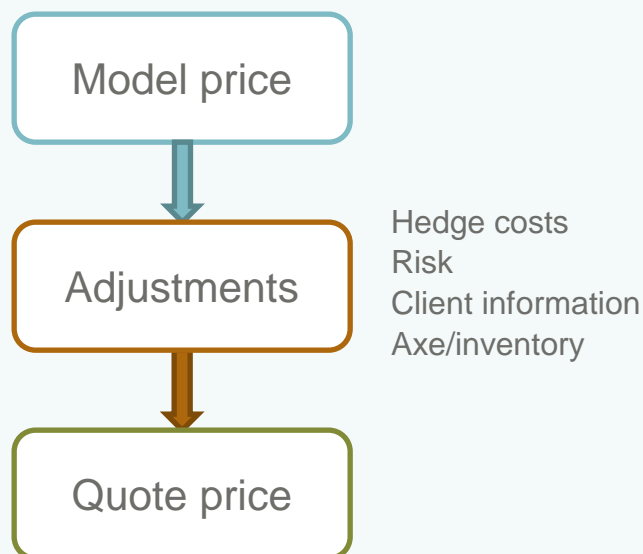
- Provide prices based on replication arguments

$$V_t = \mathbb{E}^Q[V_T]$$

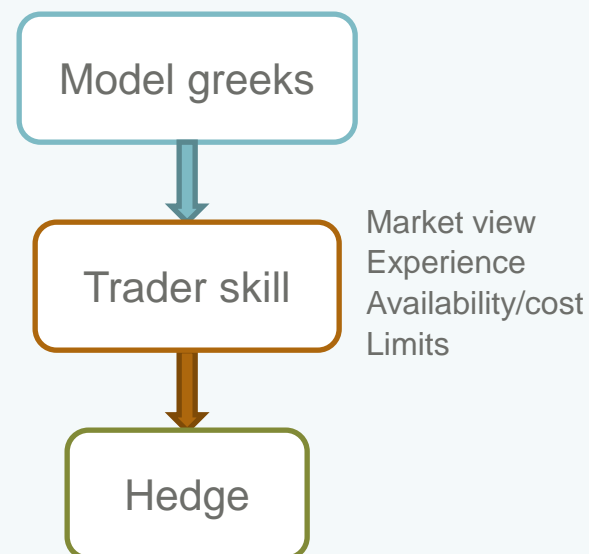
- For simple exchange-traded products, the inputs to the pricing model are fitted to the market
- The model is used to interpolate the price – for complex products, there is uncertainty
- Provide greeks: sensitivities to market data inputs

$$\Delta = \frac{\partial V}{\partial S}$$

Quoting



Hedging



Why is hedging hard?

Key decisions

- When to hedge
- What to buy/sell

Should I re-hedge now?

How much unhedged exposure do I have?

Am I near my risk limits?

Is my current position carrying well?

How much will it cost to hedge?

What should I buy?

Flattening all greeks is not practical

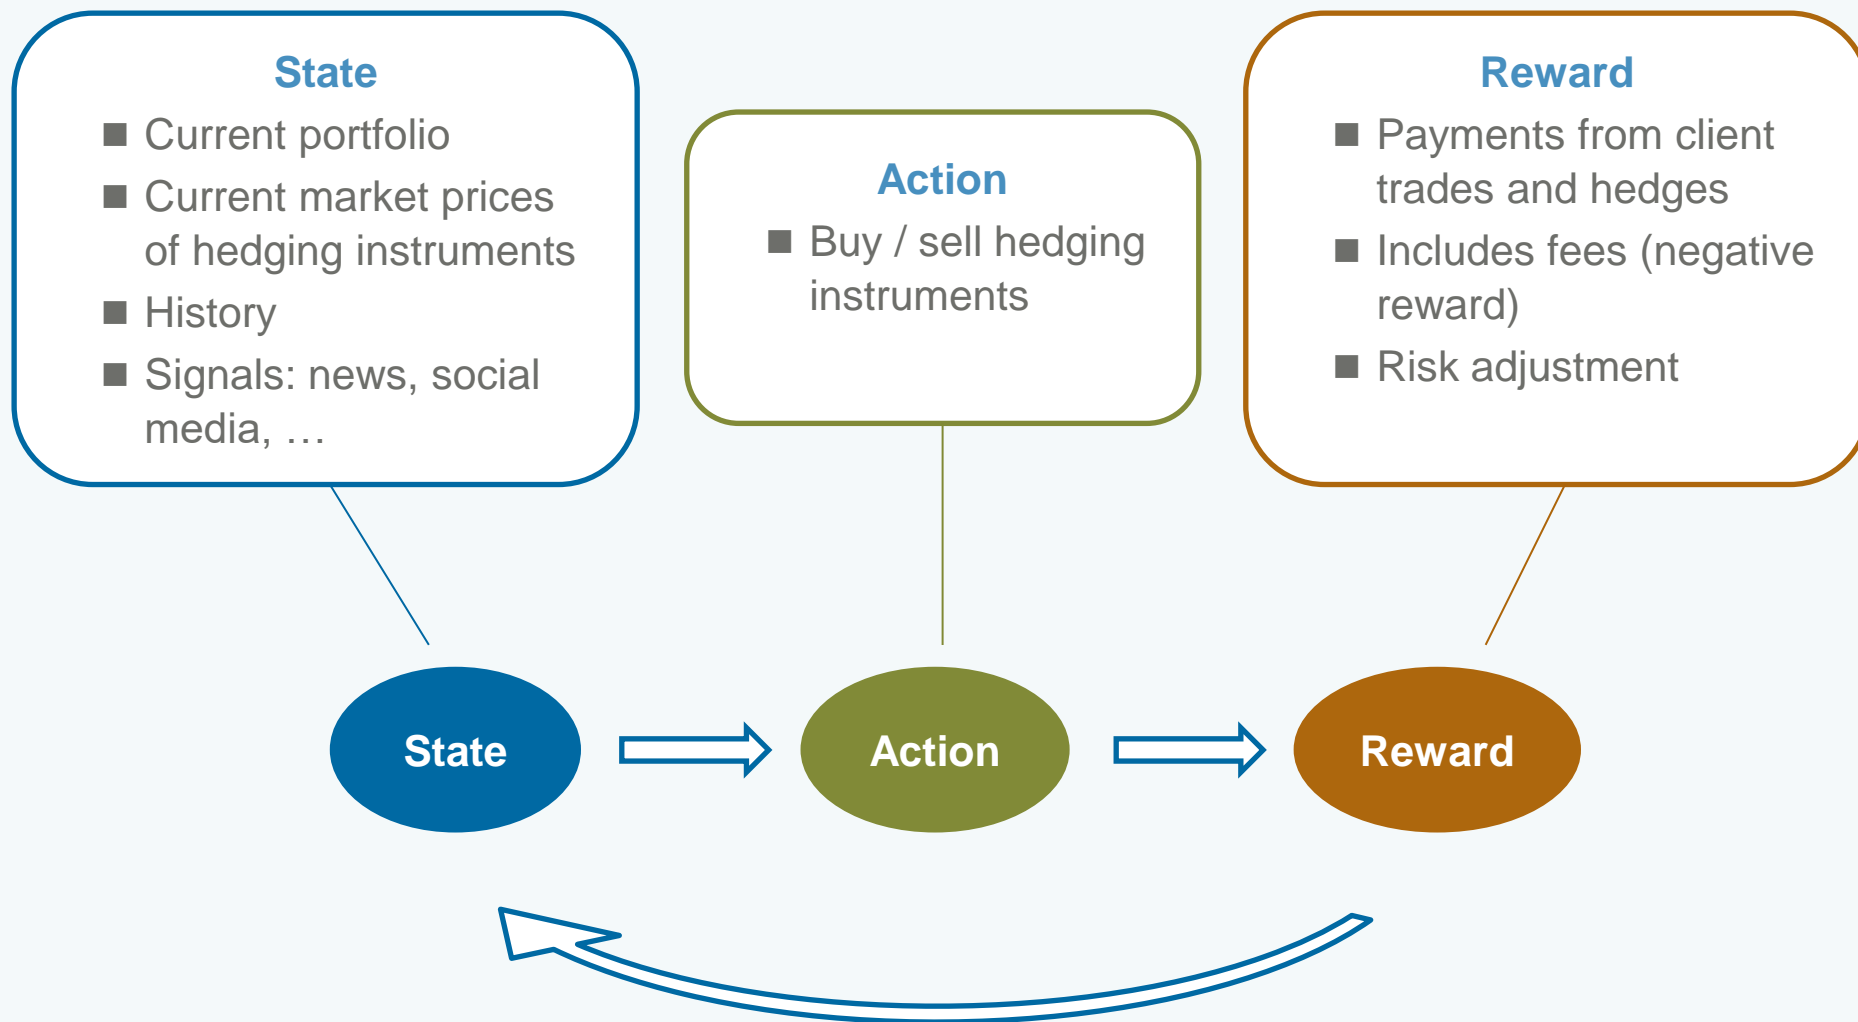
Which hedges are cheap / expensive?

How do the prices of the hedges move together?

Which hedges will need adjusting again later?

Trade-off of risk vs cost

The Reinforcement Learning paradigm is a good fit for trading



Deep Hedging principles

Problem statement

- Use AI to find optimal hedging strategies for derivatives
- Allow for important real-world effects (costs, discrete hedging, limits)
- Take a more systematic approach to hedging: less art, more science

Core ideas

- A hedging strategy is a **policy**: a function mapping state to action
- **State** includes the market and our portfolio
- **Actions** involve buying or selling liquid hedging instruments
- Cash payments from buying, selling, or from our existing portfolio are our **rewards**
- The policy will determine the profit and loss from hedging on any future path
- We define a **loss function** on the distribution of hedged P&L
- We optimize the policy with respect to this objective

Episodic formulation

- Learn to hedge a **specific portfolio** Z to maturity

- Write the terminal gain of a set of hedging actions:

$$G^Z(a) = Z_T + \sum_{t=0}^{m-1} (\delta_t \cdot (H_{t+1} - H_t) - c_t(a_t)) = Z_T + \sum_{t=0}^{m-1} (a_t \cdot (H_T - H_t) - c_t(a_t))$$

- Model the **action** as the output of a neural network, which represents our **policy**

$$a_t = a^\pi(s, t; \theta)$$

- Maximize the **utility** of the terminal gain distribution

$$\mathcal{L} = -U(G^Z(\theta))$$

- Obtain sample paths of the market state, and evaluate the portfolio payments on each path

$$(H_0, H_1, \dots, H_T)^i \rightarrow Z_T^i$$

- Train by applying **stochastic gradient descent** to the loss function in batches of samples

$$\theta \rightarrow \theta - \gamma \sum_i \nabla_\theta \mathcal{L}_i$$

Deep Hedging gains

Deeper dive into the gains process

- Write the terminal gain of a set of hedging actions:

$$G^Z(a) = Z_T + \sum_{t=0}^{m-1} (\delta_t \cdot (H_{t+1} - H_t) - c_t(a_t)) = Z_T + \sum_{t=0}^{m-1} (a_t \cdot (H_T - H_t) - c_t(a_t))$$

Portfolio cashflows

- Market state dependent
- Independent of actions

Hedge instruments

- H_t is the vector of mid prices of the available hedge instruments
- Independent of actions

Transaction costs

- Typically a convex function of the action, e.g. proportional

Actions

- a_t is the action at step t
- It depends on our policy
- Related to our hedge instrument holdings: $\delta_t = \delta_{t-1} + a_t$


Deep Hedging actions

Meaning of the policy

- The actions represent how much of each hedge instrument to buy or sell at each step, in each state
- Modelled as the output of a neural network

$$a_t = a^\pi(s, t; \theta)$$

- At each time step:



- Dependence on previous actions introduces recursion


Deep Hedging actions

Architecture

- Different choices are possible, but all reflect the recursive nature of the problem


Sequential

$$a_t = f(H_t, \delta_{t-1}; \theta_t)$$



Shared weights


$$a_t = f(H_t, \delta_{t-1}, t; \theta)$$



LSTM

$$a_t = f(H_t, \delta_{t-1}, l_{t-1}; \theta)$$

$$l_t = g(H_t, \delta_{t-1}, l_{t-1}; \theta)$$



Utility functions

What are the considerations for a utility function?

- Plays a critical role in determining the optimal policy
- Should reflect risk aversion and preference for positive PnL
- A classic choice in finance is **mean-variance**

$$U_{\lambda}(X) = \mathbb{E}[X] - \frac{\lambda}{2} \text{Var}[X]$$

λ is a risk aversion parameter

- This is okay if X is normally distributed
- We can think of λ as the cash price of a unit of variance risk

When does mean-variance go wrong?

- Mean-variance is not monotonic
 - Consider two strategies X and Y
 - If $X > Y$ in all possible outcomes, X is clearly better
 - However, we may have $\mathbb{E}[X] - \frac{\lambda}{2} \text{Var}[X] < \mathbb{E}[Y] - \frac{\lambda}{2} \text{Var}[Y]$ and mean-variance then prefers Y

Utility functions

Example

- Add a **free** lottery ticket to a portfolio

$$X = Z$$

$$Y = Z + L$$

L is the lottery ticket

$$\begin{aligned} U^\lambda(Y) &= \mathbb{E}[Y] - \frac{\lambda}{2} \text{Var}[Y] = U^\lambda(X) + \mathbb{E}[L] - \frac{\lambda}{2} \text{Var}[L] \\ &= U^\lambda(X) + pN \left(1 - \frac{\lambda}{2} (1-p)N \right) \end{aligned}$$

- The free ticket is **rejected** if $\lambda > \frac{2}{N(1-p)}$

Desirable properties of utility functions

- Monotonicity

$$X \geq Y \Rightarrow U(X) \geq U(Y)$$

- Concavity, because we are risk-averse

$$U(\alpha X + (1 - \alpha)Y) \geq \alpha U(X) + (1 - \alpha)U(Y)$$


- Cash-invariance, to give U the meaning of a value

$$U(X + c) = U(X) + c$$

Training data

Requirements

- We need samples of paths of the market, as long as the lifetime of the portfolio to be hedged
- For proof of concept, we can generate synthetic data
 - Use simple, off-the-shelf classical models, e.g. Black-Scholes, Heston, local volatility
 - Has the advantage of a baseline for performance
- For production use, we need more realistic data



Building the training dataset

Generate paths of hedge instruments

- Include asset spot prices, but usually also a grid of vanilla options
- Typically daily sampled
- Hedge instrument prices, including transaction costs, and payoffs



Decorate each path with cashflows from the portfolio to be hedged

- These are independent of actions

Training

Optimization approach

- Finite time horizon
- Continuous, high-dimensional state and action space
- State is largely independent of actions: $S = (M, Z, \delta)$
- Objective based on terminal utility
- All motivate the choice of **gradient-based direct policy search**
 - Not common in the RL community
 - Related to REINFORCE, but deterministic policy

Stochastic gradient descent

- Compute the loss function for the current policy on a batch of paths
- Update the network parameters by following the gradient

$$\theta \rightarrow \theta - \gamma \sum_i \nabla_{\theta} \mathcal{L}_i$$

- Vanilla SGD / Adam / RMSProp

Episodic formulation

- Learn to hedge a **specific portfolio** Z to maturity

- Write the terminal gain of a set of hedging actions:

$$G^Z(a) = Z_T + \sum_{t=0}^{m-1} (\delta_t \cdot (H_{t+1} - H_t) - c_t(a_t)) = Z_T + \sum_{t=0}^{m-1} (a_t \cdot (H_T - H_t) - c_t(a_t))$$

- Model the **action** as the output of a neural network which represents our **policy**

$$a_t = a^\pi(s, t; \theta)$$

- Maximize the **utility** of the terminal gain distribution

$$\mathcal{L} = -U(G^Z(\theta))$$

- Obtain sample paths of the market state, and evaluate the portfolio payments on each path

$$(H_0, H_1, \dots, H_T)^i \rightarrow Z_T^i$$


- Train by applying **stochastic gradient descent** to the loss function in batches of samples

$$\theta \rightarrow \theta - \gamma \sum_i \nabla_\theta \mathcal{L}_i$$

Toy examples


Example: Delta-hedging in a Black-Scholes world

- Delta-hedging a call option with transaction costs in a Black-Scholes world
- Compare with the known theoretical result



Example: Delta-hedging in a Heston world




- Delta-hedging a call option with transaction costs in a Heston world
- No theoretical result



Toy examples – visualization

Example: Delta-hedging a call option in a Heston world


- Compare with the risk-neutral model result



Performance metrics

Out of sample performance

- PnL distribution
- Expected costs
- Easy to generate additional data in toy model settings
- Harder in the real world



Finding ways to understand hedging behaviour

- Visualisation becomes more challenging as the number of hedge instruments increases
- Plot quantiles of actions / cross-sections of network output

Beyond toy models

- Training data is the next challenge

Moving to the real world

Do we have enough real-world data?

- Let's say we want to learn to hedge a product with 1Y maturity
 - Hedge frequency will usually be daily
- With 10Y of historical option price data, we have 10 fully independent paths
- Even if we allow overlapping paths, we only have ~2500 samples
- **Not enough data** to train a network hedger directly

What can we do?

- Create realistic synthetic data
- This means building market simulators
- For our equity derivatives applications, we need to learn to simulate the entire vanilla option market

Information content

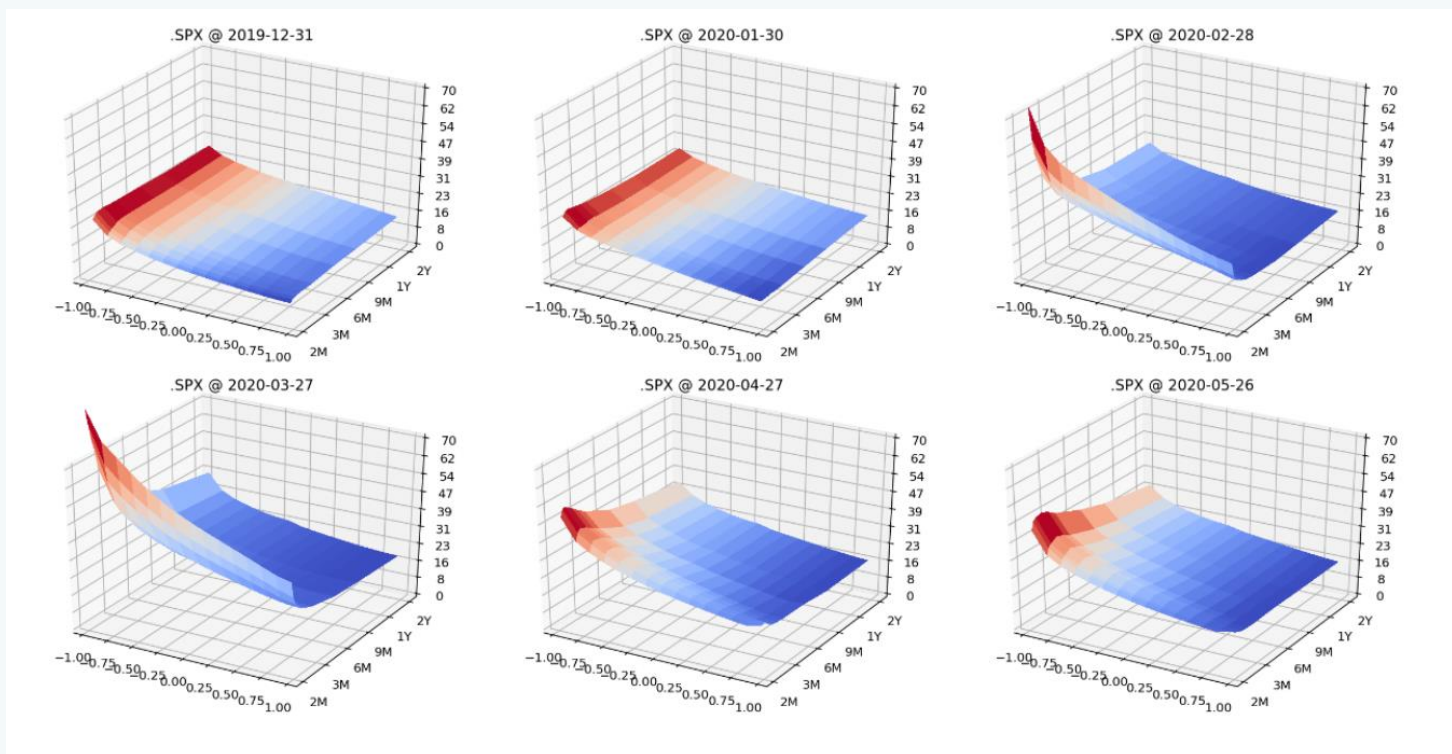
- A simulator trained on historical data adds no new information
- Instead it provides controlled, independently testable, smooth interpolation of the data
- Could be overlaid with additional features by a human expert: alpha, tail risk

What does the option market look like?

- For major indices, thousands of listed call and put options with different **strikes** and **maturities** are available and traded in volume at any time
- We will aim to generate realistic daily time series for a coarser grid, e.g. 4×9

Historical option price grids for SPX

- Option prices are conventionally described in terms of implied volatilities




Equity option markets

Volatility time series

- Skew / smile
- Clustering – high autocorrelation
- Mean reversion
- High cross-correlation of levels and returns
- Generally negative correlation of returns with spot

Time series of reparameterized volatilities for SX5E options



Market simulation

A simulator generates the next state in a realistic way

- Historical market states $(x_t)_t \sim p$
- Build a network-based simulator G_θ
$$X_{t+1} = G_\theta(Z_t; x_t, \dots, x_{t-p+1}) \sim p_\theta(\cdot; x_t, \dots, x_{t-p+1})$$
- Takes a historical state x_t and iid noise Z_t , and generates a new random state X_{t+1}
- Objective: calibrate G_θ such that p_θ is “close” to p

Measuring realism


- Use a range of performance metrics to assess simulated data against historical
 - Distributional metrics
 - Unconditional moments, density, cross-correlations
 - Dependency metrics
 - Autocorrelation of returns, levels

Challenges

- High dimension of x_t , e.g., 40
- Small dataset: ~2500 samples
- State dependence: next step is conditional on current state
- Tail sampling: by definition, tail events are rare and hard to sample

Address the high dimension of the data


- Exploit the high cross-correlation and look for a low-dimensional representation
- Any encoding should be invertible – we need to simulate the full grid
- Minimize information loss under encoding/decoding round trip
- Target a “nice” distribution of samples in latent space



Autoencoder structure



- Image loss objective:

$$\mathcal{L} = \sum_i \left(x_i - D_{\theta}(E'_{\theta}(x_i)) \right)^2$$




Autoencoder performance

- Linear compression (PCA) provides a baseline
- Network-based compression is approximately twice as efficient



- Compressed representation time series:




Time series generation

Generator structure

- Build a network-based simulator G_θ for encoded state variables y_t

$$Y_{t+1} = G_\theta(Z_t; y_t, \dots, y_{t-p+1}) \sim p_\theta(\cdot; y_t, \dots, y_{t-p+1})$$

- The generator takes the current state and a sample of iid noise, and gives the next state
- Repeatedly applying the generator allows us to build time series
- Classical approaches: VAR / GARCH
- Network-based generators can offer better performance



Time series generation

Generator training

- Requires some measure of distributional distance between real and generated paths

Network-based generators

- Generative adversarial networks (GANs)

$$\min_G \max_D \mathbb{E}_{Y_{0,p} \sim \mu} \left[\mathbb{E}_\mu [\ln(D(Y)) | Y_{0,p}] + \mathbb{E}_{\nu(G)} [\ln(1 - D(Y)) | Y_{0,p}] \right]$$

- Conditional Signature Wasserstein distance

$$c\mathcal{W}^{\text{Sig}}(p, p_\theta) = \mathbb{E}_p \left(\left\| \mathbb{E}_p [\mathcal{S}(Y_{t+1,t+q}) | Y_{t' \leq t}] - \mathbb{E}_{p_\theta} [\mathcal{S}(Y_{t+1,t+q}) | Y_{t' \leq t}] \right\|_2 \right)$$

- Normalizing flows

$$\text{KL}(p, p_\theta) = -\mathbb{E}_p [\mathbb{E}_p [\ln p(Y_{t+1} | Y_{t' \leq t}) - \ln p_\theta(Y_t | Y_{t' \leq t}) | Y_{t' \leq t}]]$$

Conditional Signature Wasserstein distance

- The **Wasserstein distance** is a measure of distance between distributions (“earth mover’s distance”)

$$\mathcal{W}(p, p_\theta) = \sup_f \mathbb{E}_p[f(Y)] - \mathbb{E}_{p_\theta}[f(Y)]$$

- To compare distributions of paths, we can use a related metric based on path **signatures**
 - The signature is a path transformation using iterated integrals, with powerful properties
 - In particular, if two processes have the same expected signature, they have the same law
- The Signature Wasserstein-1 metric is

$$\mathcal{W}^{\text{Sig}}(p, p_\theta) = \|\mathbb{E}_p[\mathcal{S}(Y)] - \mathbb{E}_{p_\theta}[\mathcal{S}(Y)]\|_2$$

- For our generator, we need the **conditional** Signature Wasserstein-1 distance

$$c\mathcal{W}^{\text{Sig}}(p, p_\theta) = \mathbb{E}_p \left(\|\mathbb{E}_p[\mathcal{S}(Y_{t+1,t+q})|Y_{t' \leq t}] - \mathbb{E}_{p_\theta}[\mathcal{S}(Y_{t+1,t+q})|Y_{t' \leq t}]\|_2 \right)$$

- We still need to estimate the two conditional expected signatures
 - We can generate Monte Carlo estimates for the latter, and use regression for the former

Normalizing flows

Generating a distribution

- The previously presented generators produce a new state sampled from the optimized conditional distribution
- If we can generate a **distribution** instead, then we can more efficiently compute a standard distributional distance like the KL-divergence
- Another desirable property for a generator is invertibility, i.e. we can compute

$$Z_t = G_{\theta}^{-1}(Y_{t+1}; y_t, \dots, y_{t-p+1})$$

- Being able to back out the driving noise is useful
 - Validate distributional assumptions
 - Introduce noise structure between two or more simulators
- A **normalizing flow** generator gives us these properties

Normalizing flows

Objective function

- We consider the series of conditional cumulative distribution functions for each element in turn

$$F_1(y_1) = \mathbb{P}[Y_1 \leq y_1]$$

$$F_k(y_k; y_1, \dots, y_{k-1}) = \mathbb{P}[Y_k \leq y_k | Y_1 = y_1, \dots, Y_{k-1} = y_{k-1}]$$

- Note that indices here refer to elements, not time
- The conditional CDFs are invertible
- We can use an efficient **linear neural spline** representation, which allows us to compute $p_\theta(Y_t | \mathcal{F}_t)$
- To train the generator, we can use the expected conditional KL divergence, which we can now evaluate using a simple Monte Carlo estimate



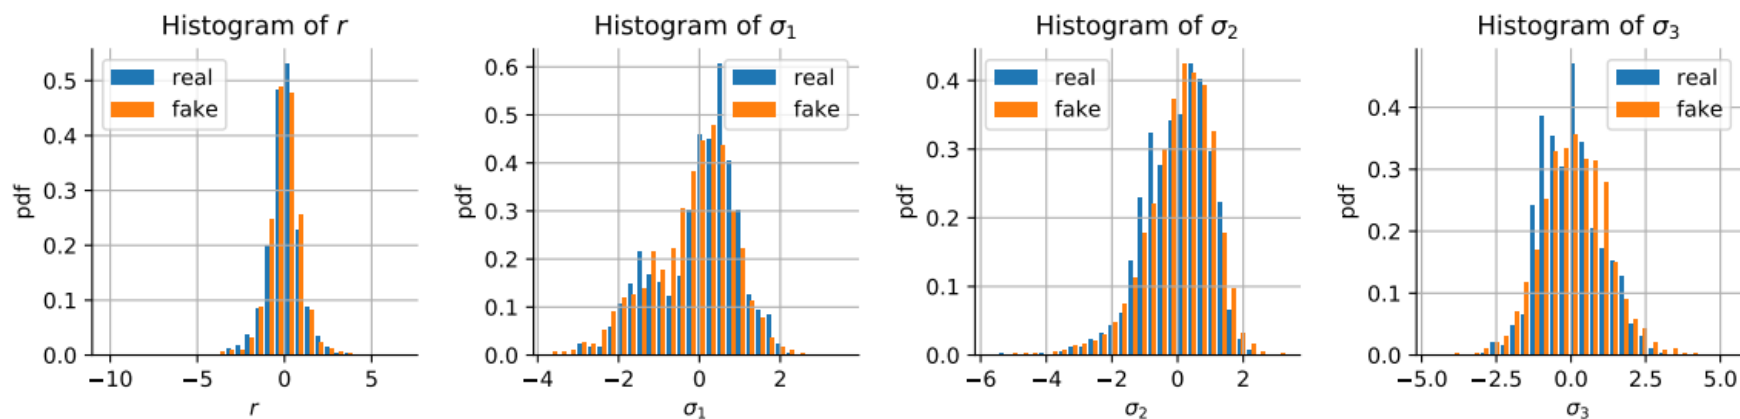
$$\text{KL}(p, p_\theta) = -\mathbb{E}_p \left[\mathbb{E}_p [\ln p_\theta(Y_{t+1} | Y_t) | Y_t] \right] + \text{const} \approx -\sum_t^T \ln p_\theta(y_{t+1} | y_t) + \text{const}$$

Challenge


- As the dimension increases, the later conditional distribution functions become noisier and harder to estimate

Normalizing flow simulator performance

Level process metrics





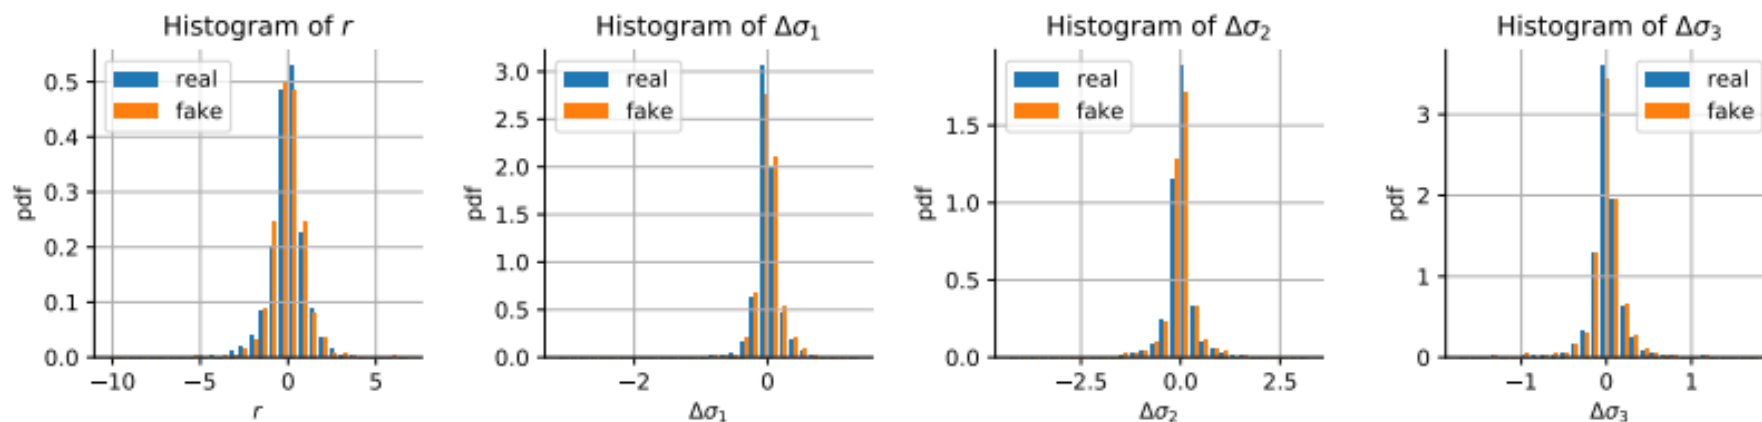
Historical cross-correlation matrix




Generated cross-correlation matrix

Normalizing flow simulator performance

Returns process metrics




Historical cross-correlation matrix



Generated cross-correlation matrix

Normalizing flow simulator performance

Joint return distributions



Extrapolation problem

- We consistently find that a small but non-zero proportion of generated paths need to be rejected: this is a consequence of **extrapolation**
- Controlling this behaviour is surprisingly difficult

Simulator research

- Multi-asset simulation, asset classes beyond equities, realized volatility, tail sampling, arbitrage control

Beyond policy search

Value function approach

- More familiar to Reinforcement Learning practitioners
- Necessary if we want to go beyond managing a fixed portfolio to expiry
- Naturally leads towards **universal** hedging

Bellman equation for Deep Hedging

- Most RL problems involve simple maximization of expected rewards, without risk aversion
 - In this setting, we can easily write the value function in terms of a **Bellman equation**

$$V^\pi(s) = \mathbb{E}[G^\pi|s] = \mathbb{E}[R^\pi(s) + V^\pi(s')|s] = R^\pi(s) + \mathbb{E}[V^\pi(s')|s]$$

- But hedging requires risk-aversion:

$$V^\pi(s) = U(G^\pi|s)$$

- Most reasonable utility functions are not **time consistent**

$$U(G^\pi) \neq U(U(G^\pi|s'')), \quad U \neq \text{exponential utility}$$

- However, exponential utility does have this property, which leads to a Bellman equation

$$V^\pi(s) = U(R^\pi(s) + G^\pi(s')) = U(R^\pi(s) + V^\pi(s')), \quad U = \text{exponential utility}$$

- The optimal policy maximizes the utility

$$V^{\pi^*}(s) = \sup_{\pi} U(R^\pi(s) + V^\pi(s'))$$

Implementing Bellman hedging

Actor / critic

- We need to train two networks
 - The **actor** learns the optimal policy
 - The loss function is related to the negative utility

$$\mathcal{L}^A = \mathbb{E} \left[e^{-\lambda (R^{\pi}(s) + V^{\pi}(s') - V^{\pi}(s))} \right]$$

- The **critic** learns the value function
 - The loss function is related to the **temporal difference** error

$$\mathcal{L}^C = \mathbb{E} \left[e^{-\lambda (R^{\pi}(s) + \bar{V}^{\pi}(s') - V^{\pi}(s))} - \lambda V^{\pi}(s) \right]$$

- Alternately update actor and critic
 - Apply averaging to stabilize critic updates

State dependence

- Both actor and critic depend on market and portfolio state
- We will use a portfolio state representation that mimics human trading

Hedging like a trader


Towards universal hedging

- Learn the optimal hedge for a general class of options and a range of risk aversion levels
- Provide the agent with prices and greeks from a simple, **wrong** model
- The agent must learn to correct the greeks and apply a suitable risk adjustment on top

Toy example

- In a world with stochastic volatility, learn to delta-hedge a class of vanilla options using Black-Scholes greeks

Impact of risk aversion on hedged PnL




Details


- Heston model world
- Provide Black-Scholes greeks as features
- Train on fixed-maturity, fixed-notional options with different strikes (one option at a time)
- Hedge with spot only
- Proportional transaction costs
- Finite time horizon (1M)
- Note: no statistical arbitrage

Impact of risk aversion on values


- The risk adjustment of an option position increases with risk aversion
- Converge to the risk-neutral model price as $\lambda \rightarrow 0$



Relative option value adjustment as a function of option strike and risk aversion



Option value adjustment as a function of strike, for different risk aversion levels




Critic performance: predicted vs measured utility


General risk-adjusted hedging

Impact of risk aversion on actions

- Hedging activity smoothly increases with risk aversion
- If we change our risk aversion mid-strategy, the agent adapts



Changing risk aversion



Initial hedge as a function of option strike and risk aversion

Conclusion

Deep Hedging

- Formulate hedging a derivatives portfolio as a reinforcement learning problem
- Use a loss function that penalizes risk
- Represent the hedging strategy as a neural network
- Solve the episodic problem with direct policy search

Market simulation

- Realistic synthetic data is necessary to train a Deep Hedging agent effectively
- Networks allow efficient encoding/decoding and realistic time series generation

Bellman hedging

- Points the way to more general, more powerful Deep Hedging agents
- We can learn the optimal risk-adjusted hedge from the greeks of a simple, wrong model

Disclaimer

Opinions and estimates constitute our judgement as of the date of this Material, are for informational purposes only and are subject to change without notice. It is not a research report and is not intended as such. Past performance is not indicative of future results. This Material is not the product of J.P. Morgan's Research Department and therefore, has not been prepared in accordance with legal requirements to promote the independence of research, including but not limited to, the prohibition on the dealing ahead of the dissemination of investment research. This Material is not intended as research, a recommendation, advice, or or solicitation for the purchase or sale of any financial product or service, or to be used in any way for evaluating the merits of participating in any transaction. Please consult your own advisors regarding legal, tax, accounting or any other aspects including suitability implications for your particular circumstances. J.P. Morgan disclaims any responsibility or liability whatsoever for the quality, accuracy or completeness of the information herein, and for any reliance on, or use of this material in any way. Important disclosures at: www.jpmorgan.com/disclosures

References

Deep Hedging

- **Deep Hedging**, Hans Buehler, Lukas Gonon, Josef Teichmann, Ben Wood
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3120710
- **Deep Hedging: Hedging Derivatives Under Generic Market Frictions Using Reinforcement Learning**, Hans Buehler, Lukas Gonon, Josef Teichmann, Ben Wood, Baranidharan Mohan, Jonathan Kochems
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3355706

Market simulation

- **Multi-Asset Spot and Option Market Simulation**, Magnus Wiese, Ben Wood, Alexandre Pachoud, Ralf Korn, Hans Buehler, Phillip Murray, Len Bai https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3980817
- **Deep Hedging: Learning to Simulate Equity Option Markets**, Magnus Wiese, Len Bai, Ben Wood, Hans Buehler
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3470756

Bellman hedging

- **Deep Hedging: Continuous Reinforcement Learning for Hedging of General Portfolios across Multiple Risk Aversions**, Phillip Murray, Ben Wood, Hans Buehler, Magnus Wiese, Mikko Pakkanen
<https://arxiv.org/abs/2207.07467?context=stat.ML>
- **Deep Bellman Hedging**, Hans Buehler, Phillip Murray, Ben Wood
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4151026

Statistical arbitrage

Definition

- A market has statistical arbitrage if it admits a positive risk-adjusted return for an initially empty portfolio

$$\sup_a U(G^0(a)) > 0$$

- Intuition: with no existing portfolio, the optimal action is not to do nothing

- For reasonable (risk-averse) utility functions:

$$\sup_a \mathbb{E}[G^0(a)] = 0 \Rightarrow \text{no statistical arbitrage}$$

- We can go further:

$$H_t - \gamma_t \leq \mathbb{E}[H_T | \mathcal{F}_t] \leq H_t + \gamma_t \Leftrightarrow \text{no statistical arbitrage}$$

- Here γ_t is the small-order-size limit of trading cost per unit price

Technical conditions

- Convex trading costs
 - Proportional in the small-trade limit
- Convex, bounded admissible actions
- Bounded tradable instrument values H_t

Utility functions

- Focus on exponential utility / entropic risk

$$U_\lambda(X) = -\frac{1}{\lambda} \log \mathbb{E}[e^{-\lambda X}]$$

- λ is the risk aversion

$$U_0(X) = \mathbb{E}[X]$$

Controlling statistical arbitrage

Motivation

- Am I hedging or trading for profit?
- How well can I predict future statistical arbitrage opportunities?

Removing the drift

- Find a change of measure that removes statistical arbitrage opportunities
- We want the **minimal measure change**, with the smallest distance between \mathbb{P} and \mathbb{Q}
- Remarkably, we can derive this measure directly from the optimal statistical arbitrage strategy under \mathbb{P}

$$\frac{d\mathbb{Q}^*}{d\mathbb{P}} = \frac{e^{-G^0(a^*)}}{\mathbb{E}[e^{-G^0(a^*)}]}$$

- \mathbb{Q}^* is the closest martingale measure to \mathbb{P} with respect to the relative entropy

$$H(\mathbb{Q}|\mathbb{P}) = \mathbb{E}^{\mathbb{P}} \left[\frac{d\mathbb{Q}}{d\mathbb{P}} \log \frac{d\mathbb{Q}}{d\mathbb{P}} \right]$$


- This is the **minimal entropy martingale measure**
- The result generalizes to other utility functions and to trading with transaction costs
- With transaction costs, we obtain a **near-martingale measure**

Toy example: Black-Scholes model with spot price drift

- Trade spot
- Statistical arbitrage strategy in \mathbb{P} :
buy and hold


$$\frac{dQ^*}{dP} = \exp\left(-\frac{\mu}{\sigma}W_T - \frac{\mu^2}{2\sigma^2}T\right)$$

$$H(Q^*|\mathbb{P}) = \frac{\mu^2}{2\sigma^2}T$$



Toy example: Black-Scholes model with implied volatility risk premium


- Trade spot and vanilla calls
- Statistical arbitrage strategy in \mathbb{P} :
sell options and delta hedge
- Change of measure should upweight paths with high realized volatility, and downweight paths with low realized volatility



Removing drift from a realistic option market model

VAR model for spot and option prices

- Fit a VAR(2) model to historical spot and (reparametrized) option prices for EURO STOXX 50
- Floating grid of relative-strike, relative-maturity options
- Check expected option payoffs against prices, before and after drift removal



Remark

- On a discrete measure, the boundary between static and statistical arbitrage is less clear