

Introduction to Logic Reasoning

2022 OxML Summer School

June 29, 2022

The Band

Speaker

Yitao Liang
Assistant Professor, PKU



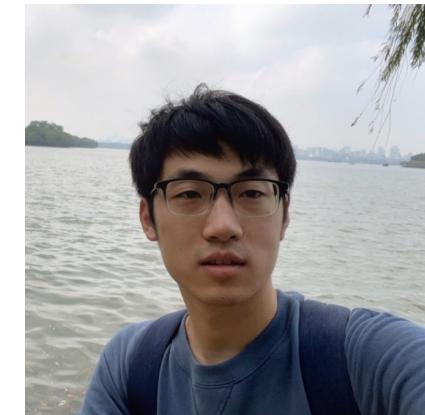
TAs



Anji Liu
UCLA



Zihao Wang
PKU



Shaofei Cai
PKU

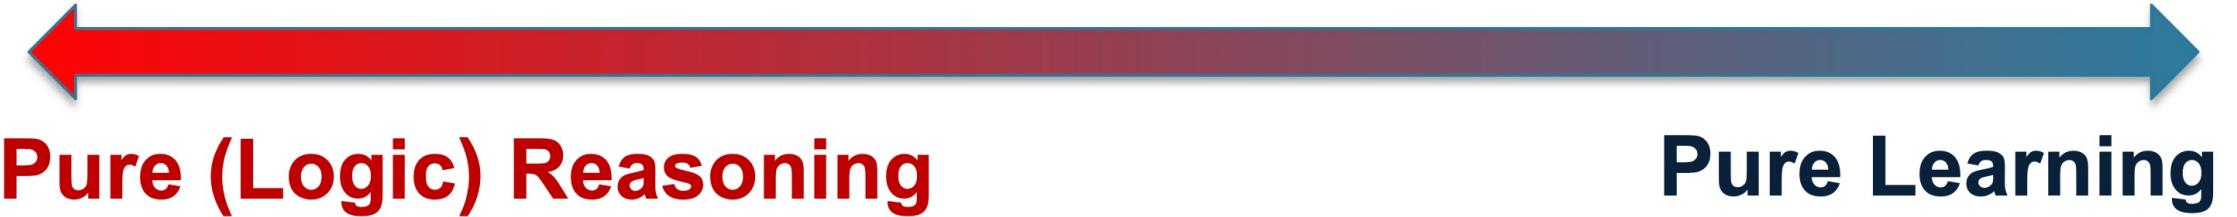
Objectives

To get you started on using logic in ML

A representation / framework

Three libraries

The AI Dilemma



The AI Dilemma



Pure (Logic) Reasoning

Pure Learning

- Fast thinking: instinctive, perceptive, model-free, interpolation
- Amazing achievements recently
- *“Pure learning is brittle”*
bias, algorithmic fairness, interpretability, explainability, adversarial attacks, unknown unknowns, calibration, verification, missing features, missing labels, data efficiency, shift in distribution, general robustness and safety fails to incorporate a sensible model of the world

The AI Dilemma

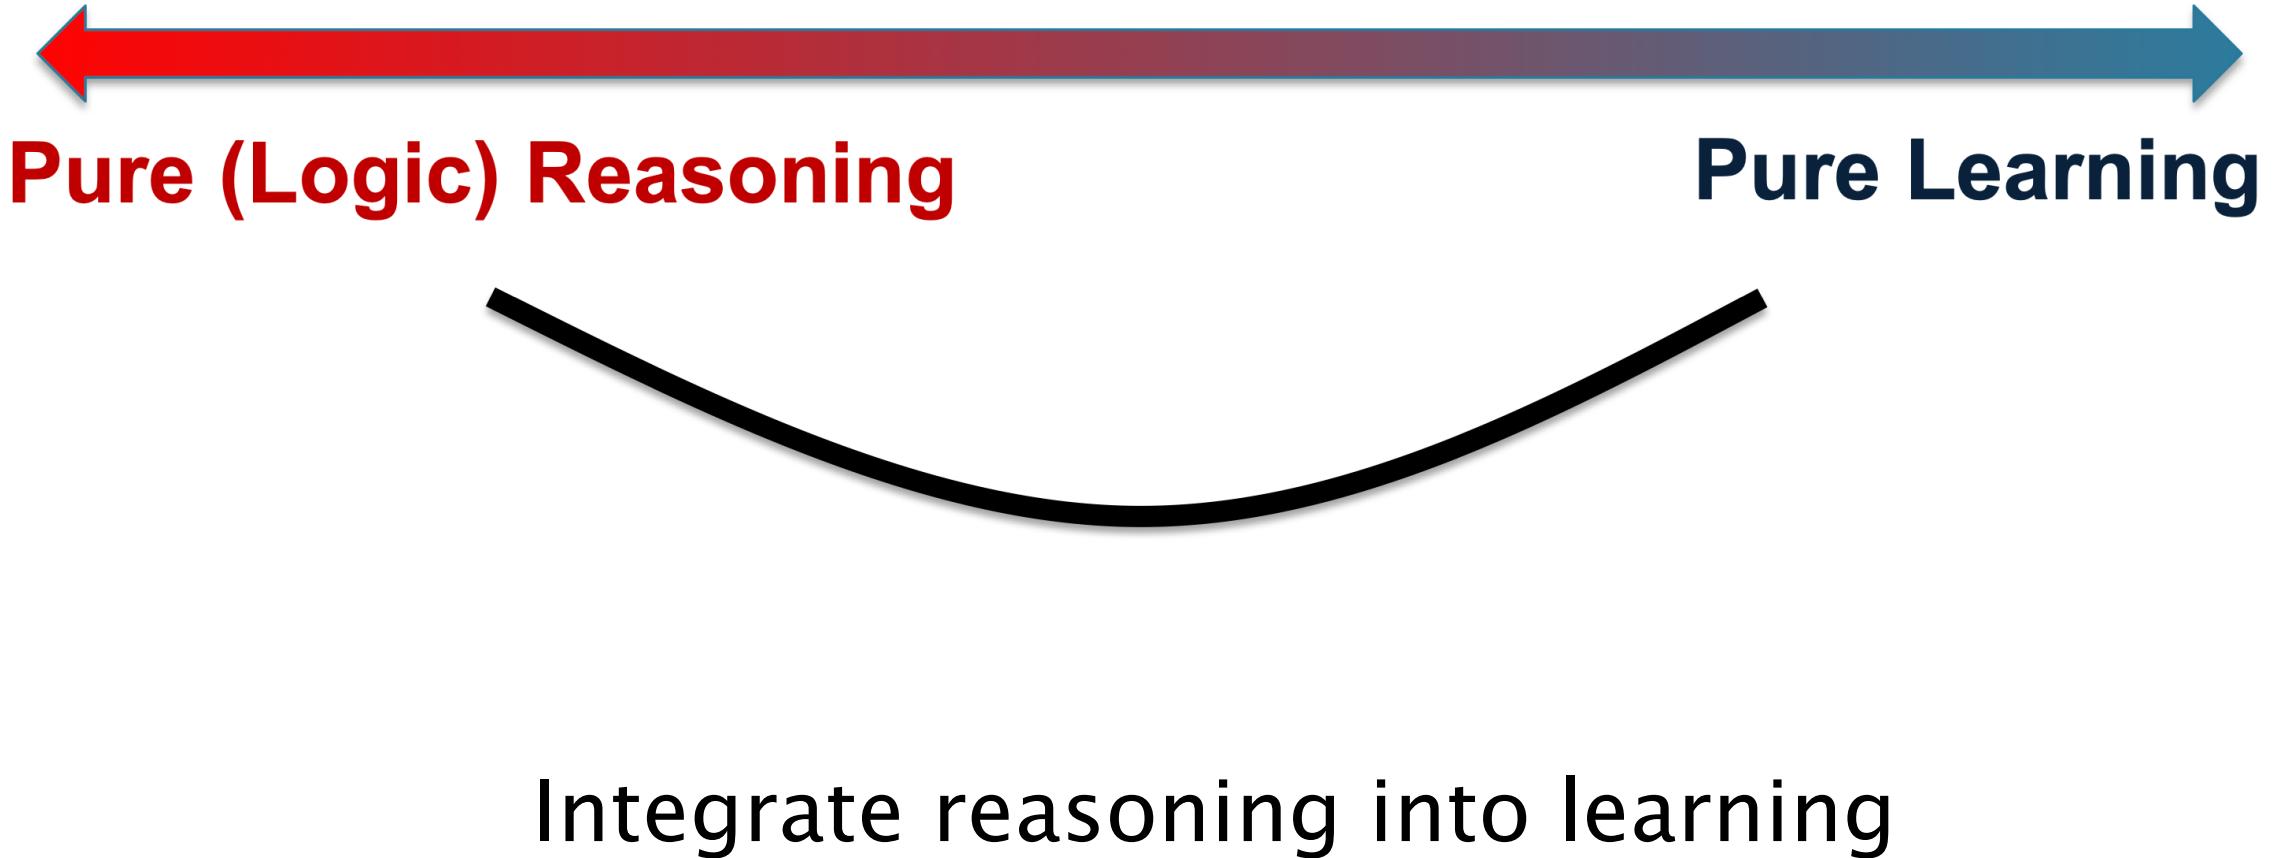


Pure (Logic) Reasoning

Pure Learning

- Slow thinking: deliberative, cognitive, model-based, extrapolation
- Amazing achievements until this day
- *“Pure logic is brittle”*
noise, uncertainty, incomplete knowledge, ...

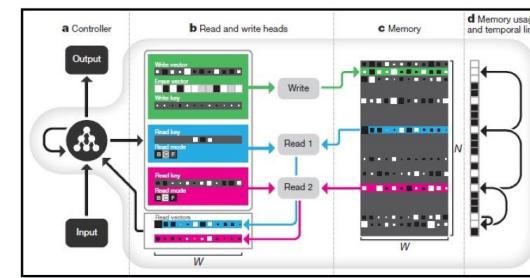
The AI Dilemma



What is the Role of Logic and Reasoning

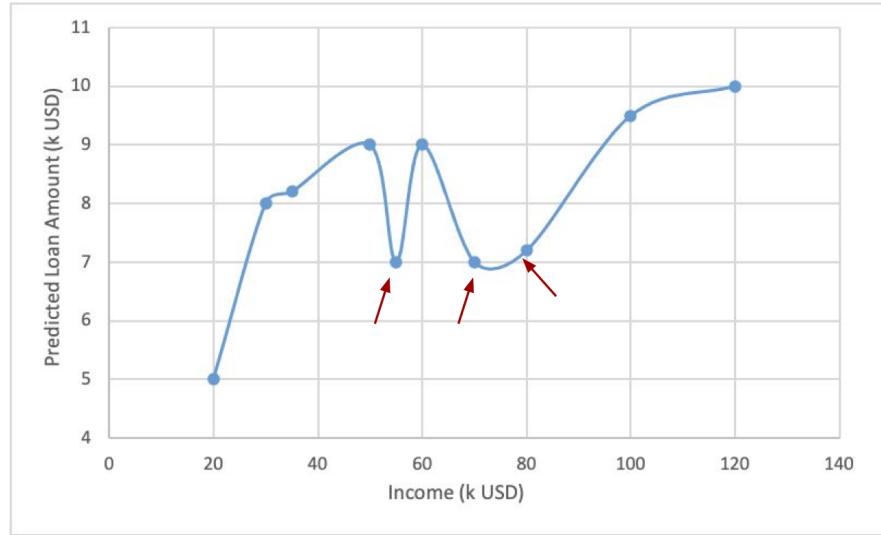
The image shows a screenshot of a New Scientist article. At the top left is the 'New Scientist' logo. Below it is a navigation bar with links: HOME, NEWS, TECHNOLOGY, SPACE, PHYSICS, HEALTH, EARTH, HUMANS, LIFE, TOPICS, EVENTS, JOBS. A green banner at the top says 'Meet The People Shaping The Future Of Energy: Reinventing Energy Summit - 25 November in London'. Below the banner is a sub-navigation bar: Home | News | Technology. There are social media sharing icons (G+, f, t, e, +) and a count of 26. The main headline reads 'DeepMind's AI has learned to navigate the Tube using memory'. Below the headline is a photograph of a person with long blonde hair looking at a detailed London Underground map.

The image shows a screenshot of a Nature article. At the top left is the 'nature' logo. Below it is a navigation bar with links: Home, News & Comment, Research, Careers & Jobs, Current Issue, Archive, Audio & Video, For Readers. Below the navigation bar is a breadcrumb trail: News & Comment > News > 2016 > November > Article. The main title is 'Google's AI reasons its way around the London Underground'. Below the title is a subtitle: 'DeepMind's latest technique uses external memory to solve tasks that require logic and reasoning — a step toward more human-like AI.' There are share and print icons at the top right.



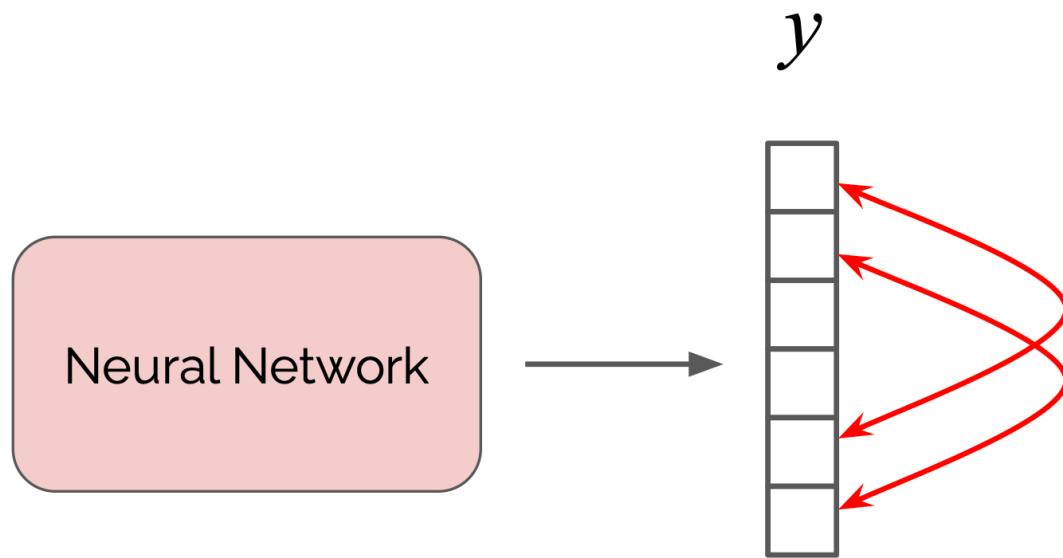
[Graves, A., Wayne, G., Reynolds, M., Harley, T., Danihelka, I., Grabska-Barwińska, A., et al.. (2016). Hybrid computing using a neural network with dynamic external memory. *Nature*, 538(7626), 471-476.]

What is the Role of Logic and Reasoning

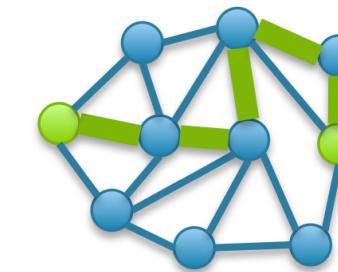


Monotonicity (Prior Knowledge):
Increasing income should increase the approved loan amount

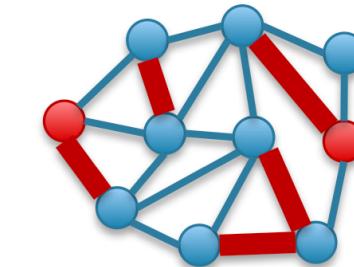
Declarative Knowledge (of the Output)



How is the output structured?
Are all possible outputs valid?



vs.



How are the outputs related to each other?

Learning this from data is inefficient
Much easier to express this declaratively

Logic Representations

Syntax

- propositional variables
- logical connectives
⇒ Boolean combinations

Syntax: Propositional Variables

Let Π be a set of *propositional variables*.

We use letters P, Q, R, S , to denote propositional variables.

Syntax: Formula

F_Π is the set of propositional formulas over Π defined inductively as follows:

F, G	$::=$	\perp	(falsum)
		\top	(verum)
		$P, \quad P \in \Pi$	(atomic formula)
		$(\neg F)$	(negation)
		$(F \wedge G)$	(conjunction)
		$(F \vee G)$	(disjunction)
		$(F \rightarrow G)$	(implication)

Syntax: Formula

F_Π is the set of propositional formulas over Π defined inductively as follows:

F, G	$::=$	\perp	(falsum)
		\top	(verum)
		$P, \quad P \in \Pi$	(atomic formula)
		$(\neg F)$	(negation)
		$(F \wedge G)$	(conjunction)
		$(F \vee G)$	(disjunction)
		$(F \rightarrow G)$	(implication)

We only need negation, And, Or

Syntax: Formula

F_Π is the set of propositional formulas over Π defined inductively as follows:

F, G	$::=$	\perp	(falsum)
		\top	(verum)
		$P, \quad P \in \Pi$	(atomic formula)
		$(\neg F)$	(negation)
		$(F \wedge G)$	(conjunction)
		$(F \vee G)$	(disjunction)
		$(F \rightarrow G)$	(implication)

How to understand implication?

SAT (Satisfiability) Reasoning

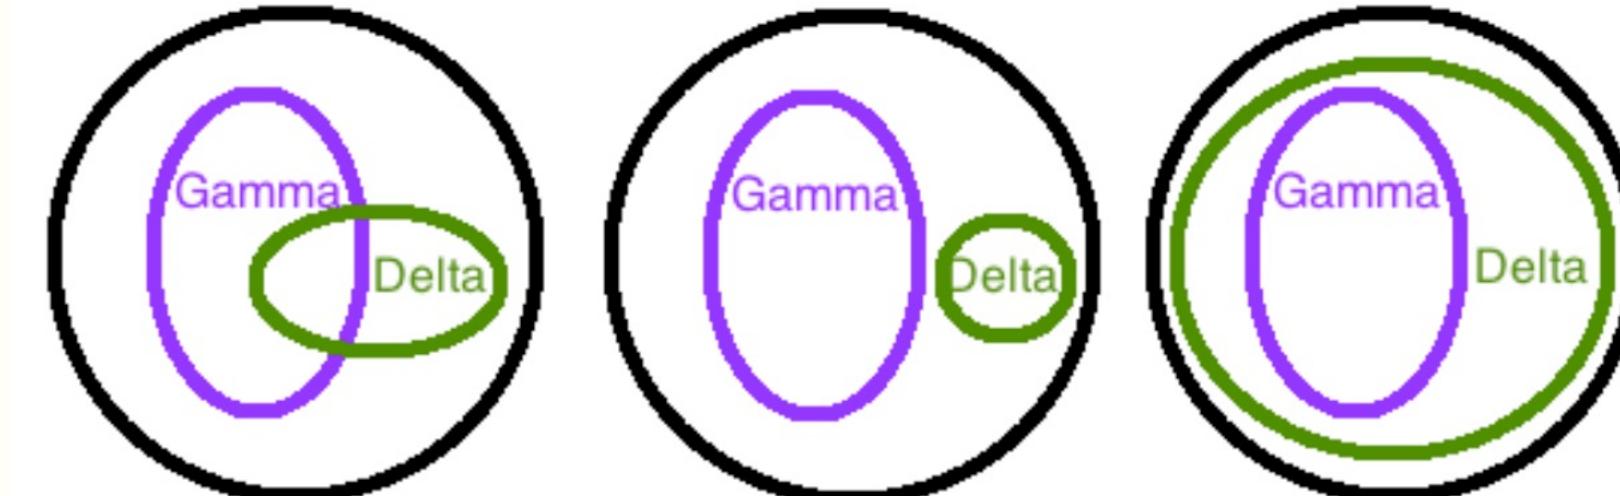
Reasoning Reduces to SAT

The basic inference problem in the propositional logic system is the following: given the knowledge that we have about the world, is it the case that some property must always hold?

In other word,

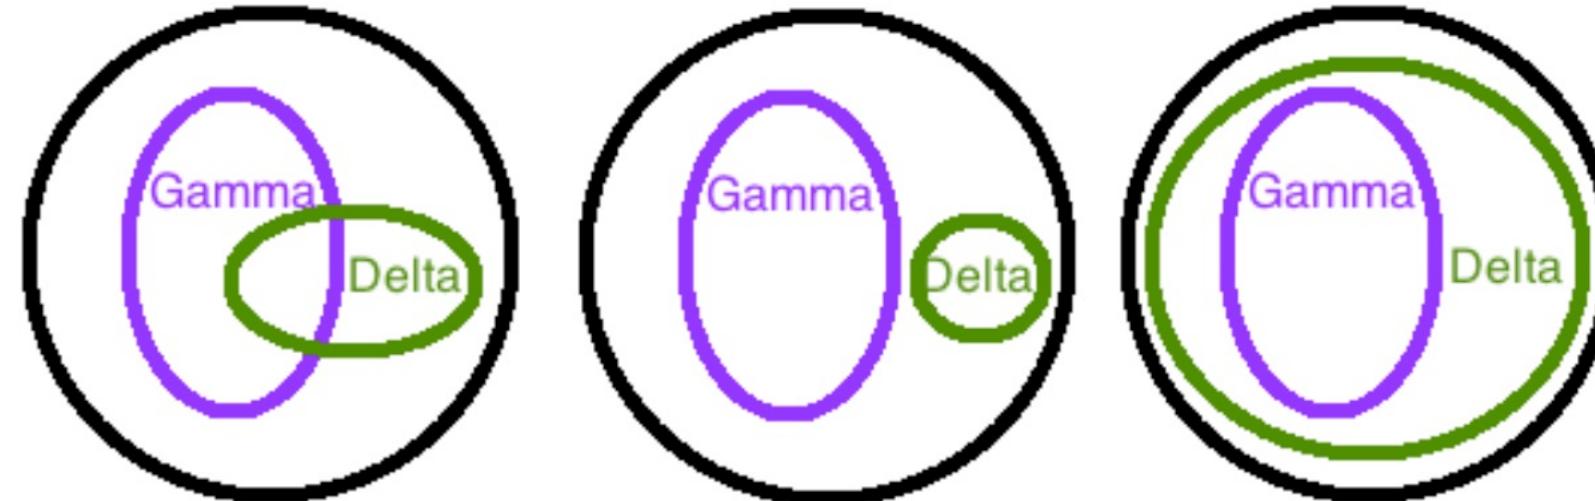
Does α imply β ?

An Exercise: SAT



After adding Delta, how will the number of possible SAT worlds change?

An Exercise: SAT



- (i) decrease; (ii) go to zero; (iii) remain the same

An Example: Sudoku

	1	2	3	4	5	6	7	8	9
1									1
2	4								
3		2							
4					5		4		7
5			8				3		
6			1		9				
7	3			4			2		
8		5		1					
9				8		6			

Goal:

Fill the empty fields with digits $1, \dots, 9$ so that each digit occurs exactly once in each row, column, and 3×3 box

An Example: Sudoku

	1	2	3	4	5	6	7	8	9
1									1
2	4								
3		2							
4					5		4		7
5			8				3		
6			1		9				
7	3			4			2		
8		5		1					
9				8		6			

Idea:

Use boolean variables $P_{i,j}^d$ with $d, i, j \in \{1, \dots, 9\}$ to encode the problem:

$P_{i,j}^d = \text{true}$ iff the value of square i, j is d

An Example: Sudoku

	1	2	3	4	5	6	7	8	9
1									1
2	4								
3		2							
4					5		4		7
5			8				3		
6			1		9				
7	3			4			2		
8		5		1					
9				8		6			

Idea:

Use boolean variables $P_{i,j}^d$ with $d, i, j \in \{1, \dots, 9\}$ to encode the problem:

$P_{i,j}^d = \text{true}$ iff the value of square i, j is d

For example:

$$P_{5,3}^8 = \text{true}$$

An Example: Sudoku

- Concrete values result in formulas $P_{i,j}^d$
- For every square (i,j) we generate $P_{i,j}^1 \vee \dots \vee P_{i,j}^9$
- For every square (i,j) and pair of values $d < d'$ we generate
 $\neg P_{i,j}^d \vee \neg P_{i,j}^{d'}$
- For every value d and row i we generate $P_{i,1}^d \vee \dots \vee P_{i,9}^d$
(Analogously for columns and 3×3 boxes)
- For every value d , row i , and pair of columns $j < j'$
we generate $\neg P_{i,j}^d \vee \neg P_{i,j'}^d$,
(Analogously for columns and 3×3 boxes)

An Example: Sudoku

Every assignment of boolean values to the variables $P_{i,j}^d$ so that all formulas become true corresponds to a Sudoku solution (and vice versa).

An Example: Sudoku

Every assignment of boolean values to the variables $P_{i,j}^d$
so that all formulas become true
corresponds to a Sudoku solution (and vice versa).

<i>world</i>	<i>A</i>	<i>B</i>	<i>C</i>
w_1	1	1	1
w_2	1	1	0
w_3	1	0	1
w_4	1	0	0
w_5	0	1	1
w_6	0	1	0
w_7	0	0	1
w_8	0	0	0

Brute force?

An Example: Sudoku

Every assignment of boolean values to the variables $P_{i,j}^d$ so that all formulas become true corresponds to a Sudoku solution (and vice versa).

<i>world</i>	<i>A</i>	<i>B</i>	<i>C</i>
w_1	1	1	1
w_2	1	1	0
w_3	1	0	1
w_4	1	0	0
w_5	0	1	1
w_6	0	1	0
w_7	0	0	1
w_8	0	0	0

Niklas Eén, Niklas Sörensson:
MiniSat (<http://minisat.se/>)

An Example: Sudoku

<i>world</i>	<i>A</i>	<i>B</i>	<i>C</i>
w_1	1	1	1
w_2	1	1	0
w_3	1	0	1
w_4	1	0	0
w_5	0	1	1
w_6	0	1	0
w_7	0	0	1
w_8	0	0	0

Niklas Eén, Niklas Sörensson:
MiniSat (<http://minisat.se/>)

Unfortunately,

The satisfiability problem is NP-complete.

Every known algorithm to solve it has an exponential time worst-case behaviour (or worse).

SAT in Practice

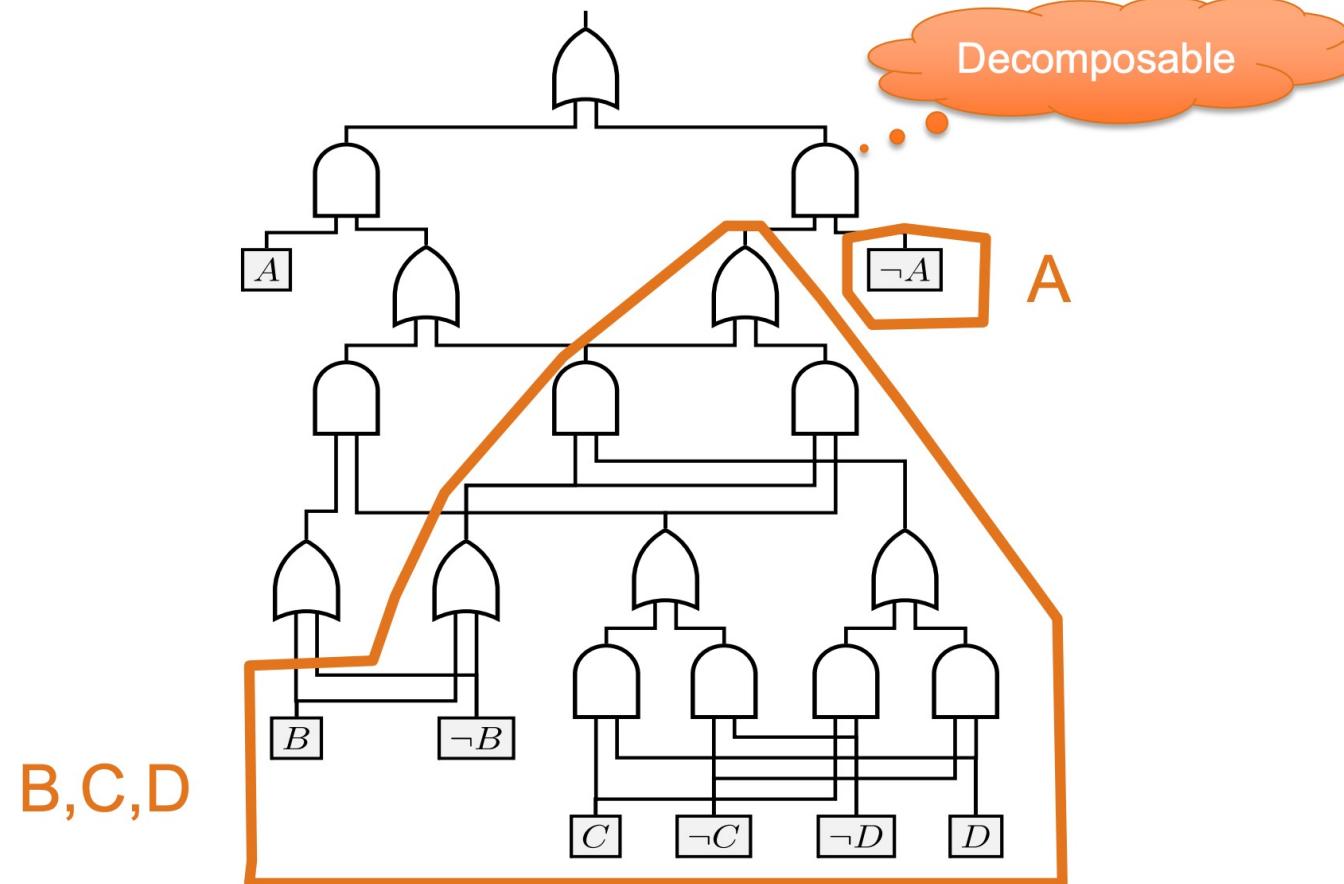
- Hardware/software verification
- Dependency checking
- Programming language compilation
- Design automation
- Combinatorial problem solvers
- Theorem proving
- ...

*Logic Representation
in Terms of SAT*

Representation for Easy SAT?

- Is there a solution? (SAT)
 - $\text{SAT}(\alpha \vee \beta)$ iff $\text{SAT}(\alpha)$ or $\text{SAT}(\beta)$ (*always*)
 - $\text{SAT}(\alpha \wedge \beta)$ iff ???

Representation: Decomposability



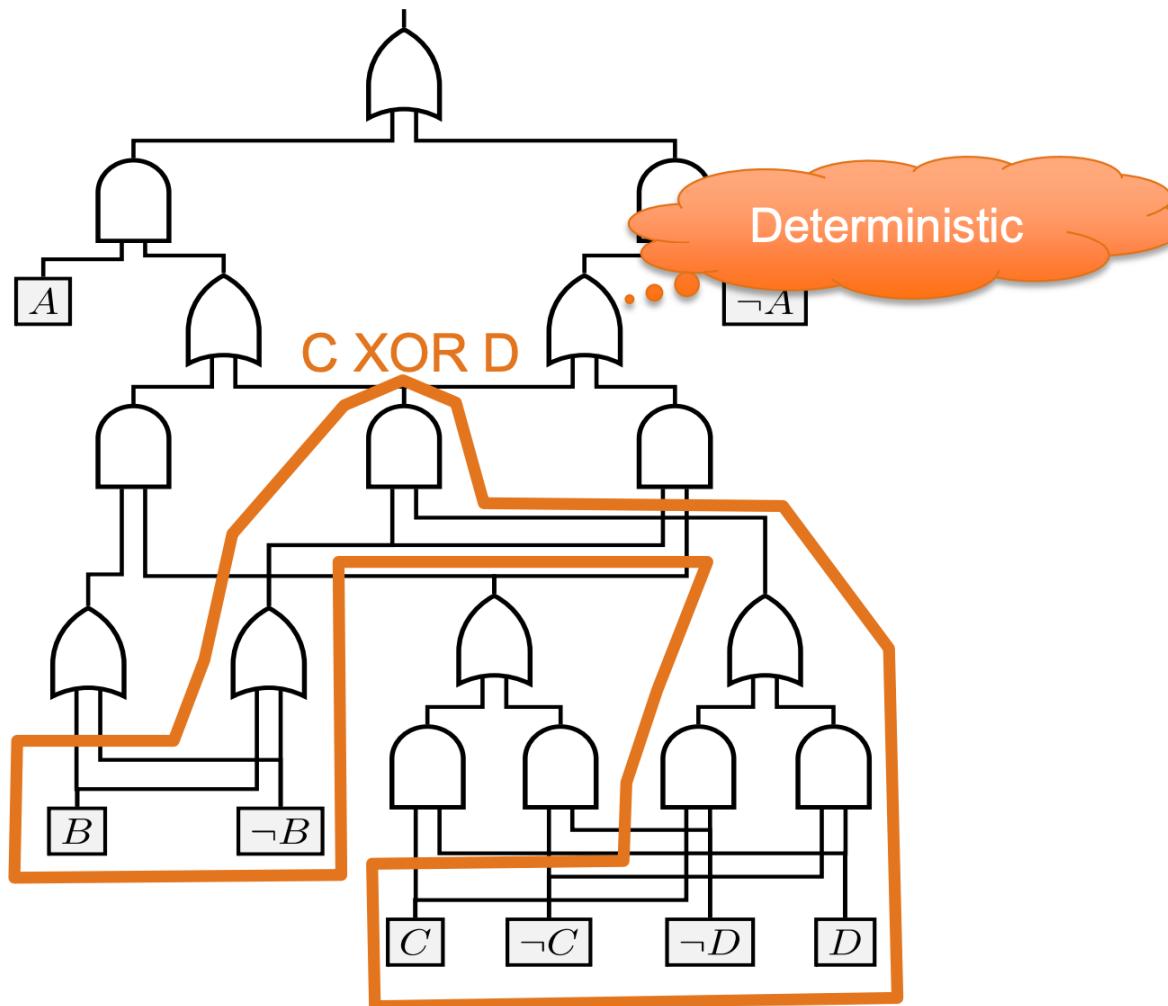
Disjoint, non-overlapping variable partition

Darwiche and Marquis, "A knowledge compilation map", 2002

Representation: #SAT

- Is there a solution? (SAT) ✓
 - $\text{SAT}(\alpha \vee \beta)$ iff $\text{SAT}(\alpha)$ or $\text{SAT}(\beta)$ (*always*)
 - $\text{SAT}(\alpha \wedge \beta)$ iff $\text{SAT}(\alpha)$ and $\text{SAT}(\beta)$ (*decomposable*)
- How many solutions are there? (#SAT)

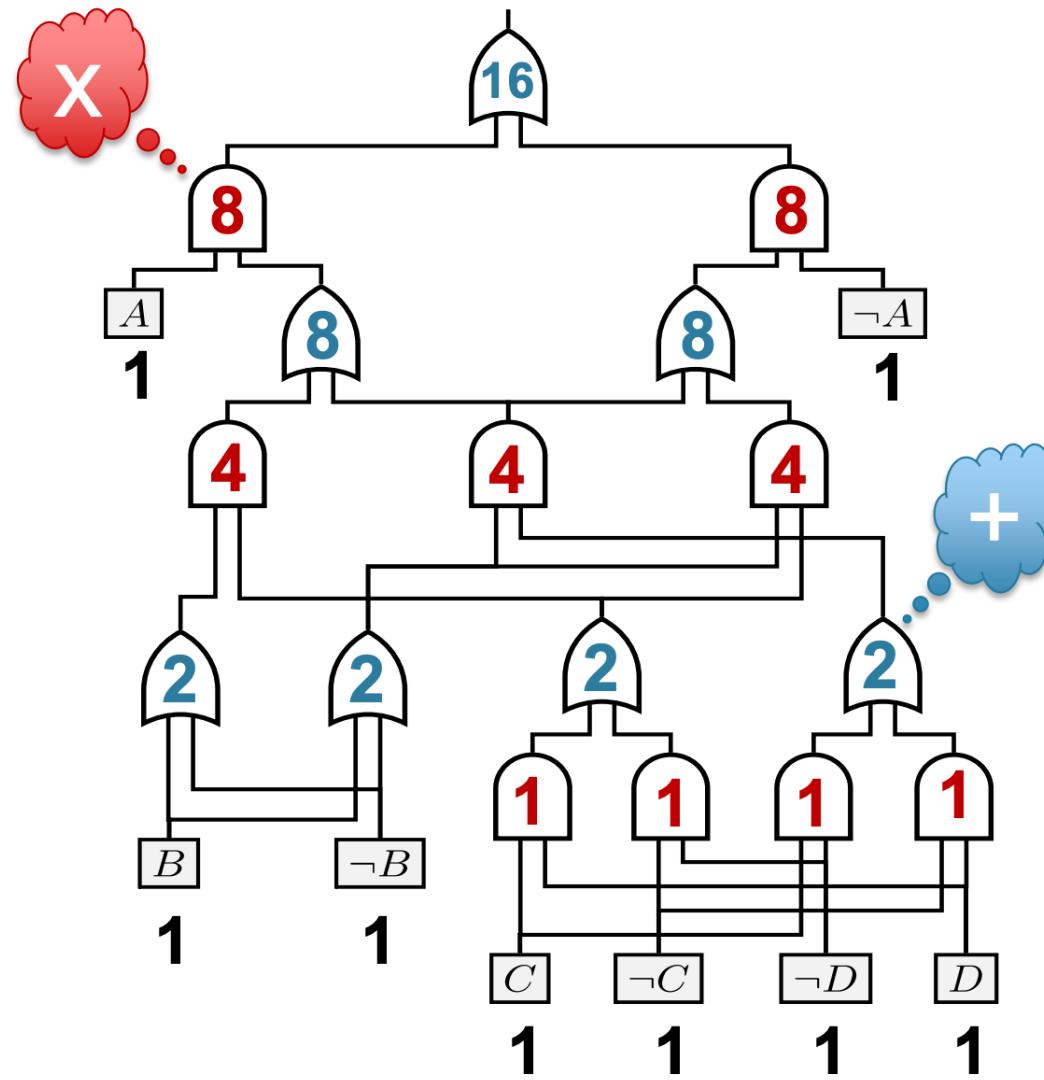
Representation: Determinism



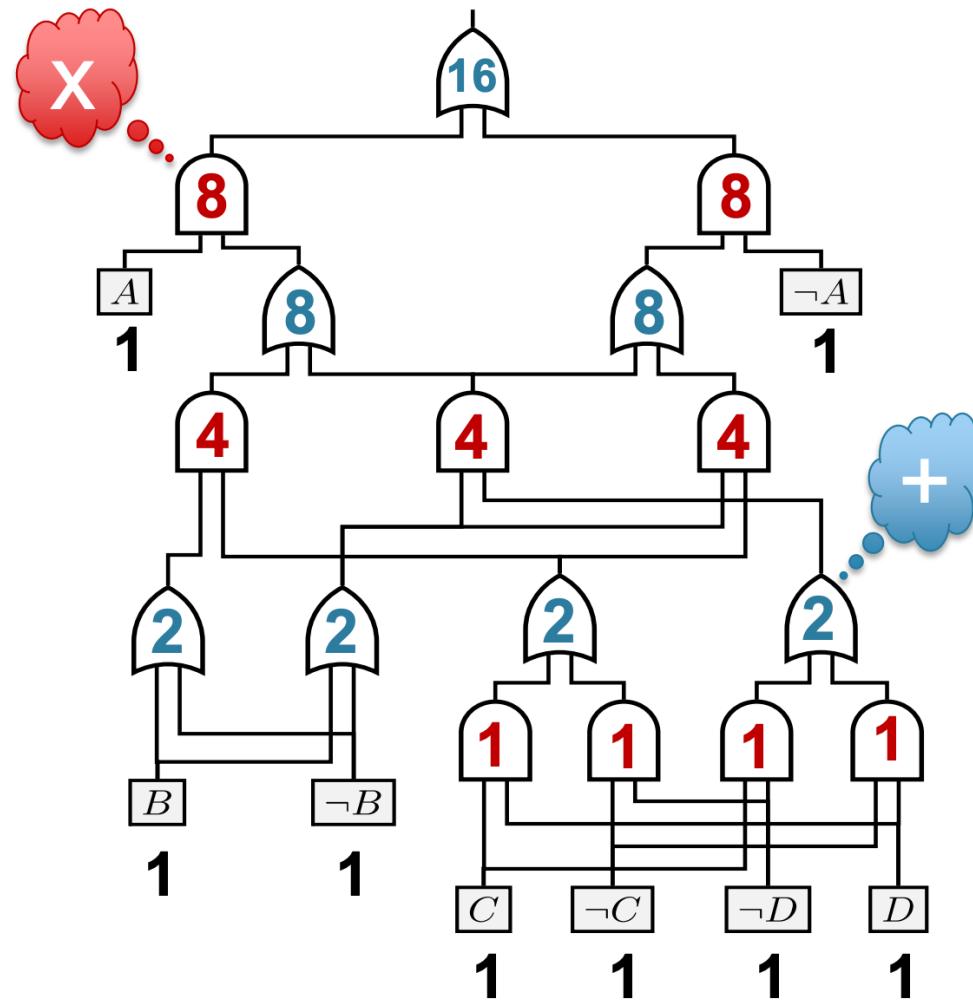
Mutually exclusive, exhaustive assignment partition

Darwiche and Marquis, "A knowledge compilation map", 2002

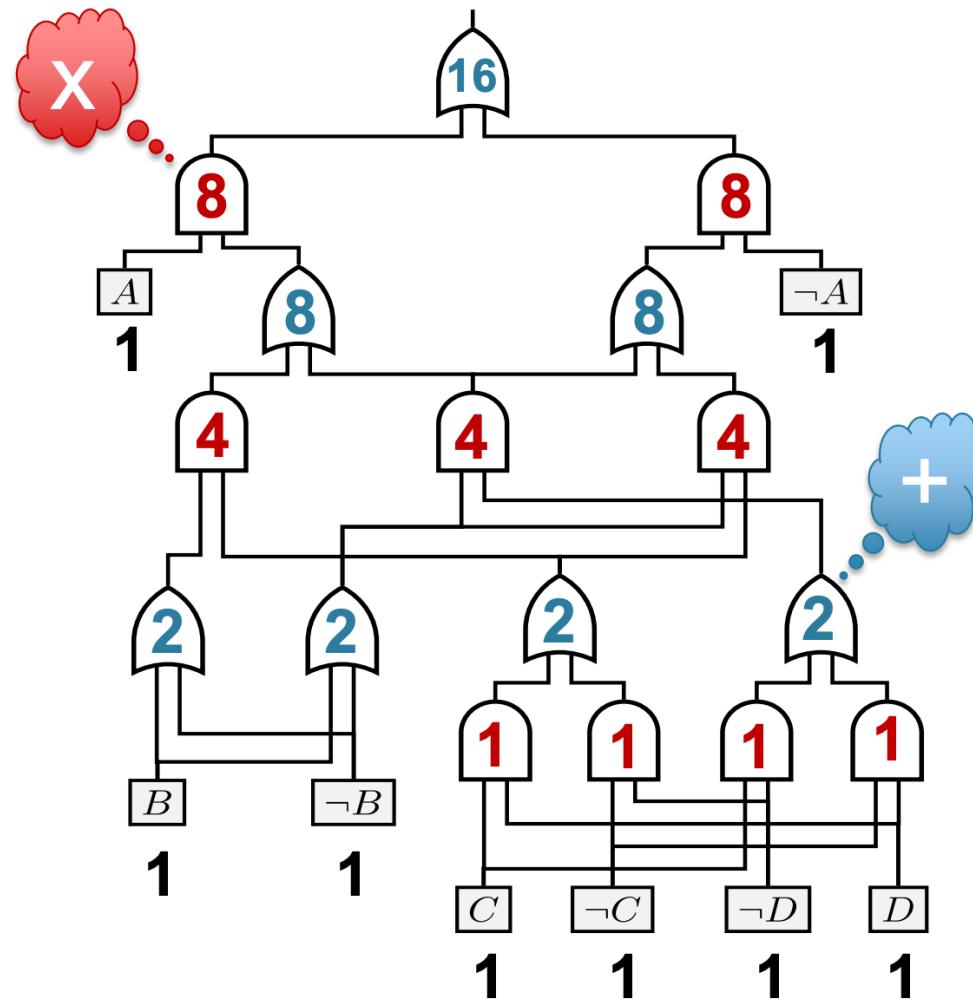
Representation: #SAT



Exercise: #SAT \rightarrow A Solution



Exercise: #SAT \rightarrow A Solution



Following a path that is non-zero

Representation: Tractable Circuits

- Is there a solution? (SAT) ✓
- How many solutions are there? (#SAT) ✓
- Complexity linear in circuit size 😊

Representation: Some Insights

- The circuit can be exponential in size of inputs
- Compiling this circuit is NP-Hard
- Then what is the benefit?

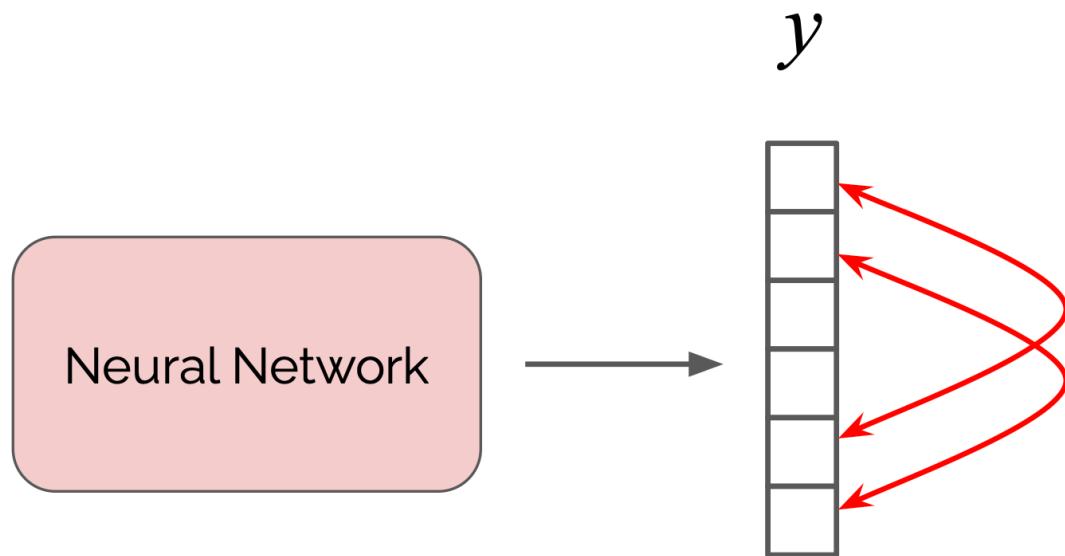
Representation: Some Insights

- The circuit can be exponential in size of inputs
- Compiling this circuit is NP-Hard
- Then what is the benefit?

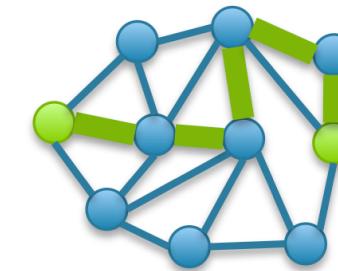
Amortization (only pay the exponential price once)

Neural Symbolic Learning

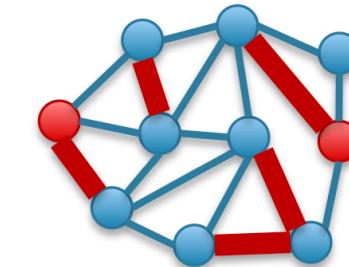
Recall: Declarative Knowledge on Outputs



How is the output structured?
Are all possible outputs valid?



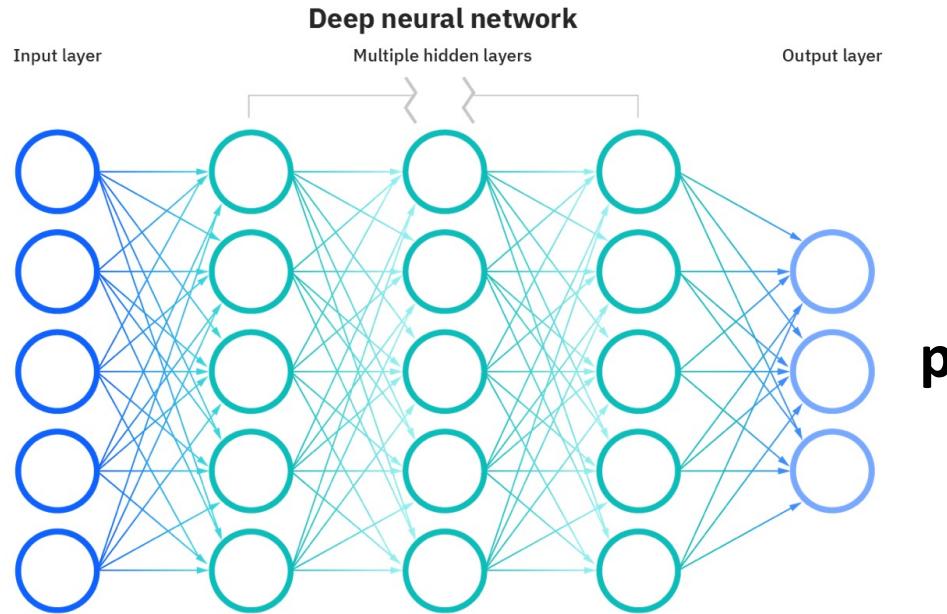
vs.



How are the outputs related to each other?

Learning this from data is inefficient
Much easier to express this declaratively

A Running Example: Classification

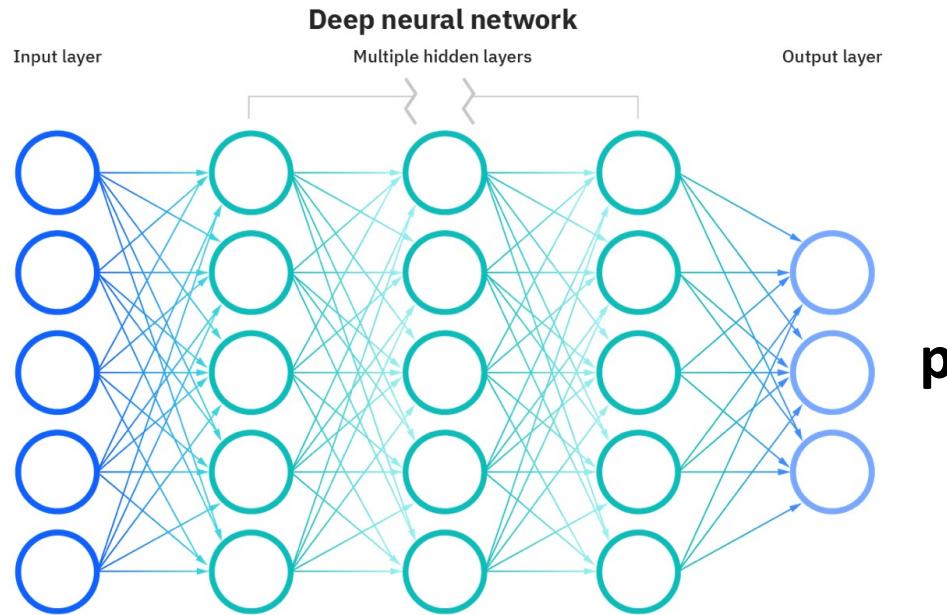


Learning subject to Exactly-One

$\neg Light \wedge Ball$
 $\neg Ball \wedge Light$
Etc.

α

A Running Example: Classification



Learning subject to Exactly-One

¬Light \wedge *Ball*
¬Ball \wedge *Light*
Etc.

α

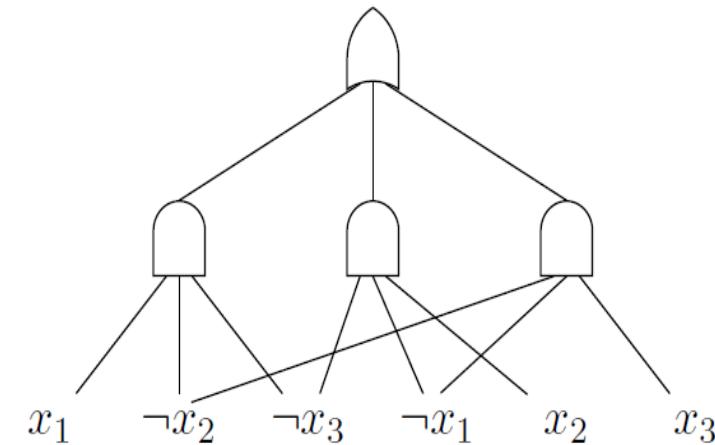
Output is
probability vector p ,
not Boolean logic!

Representation: Exactly One

Learning subject to **Exactly-One**

$\neg Light \wedge Ball$
 $\neg Ball \wedge Light$
Etc.

α



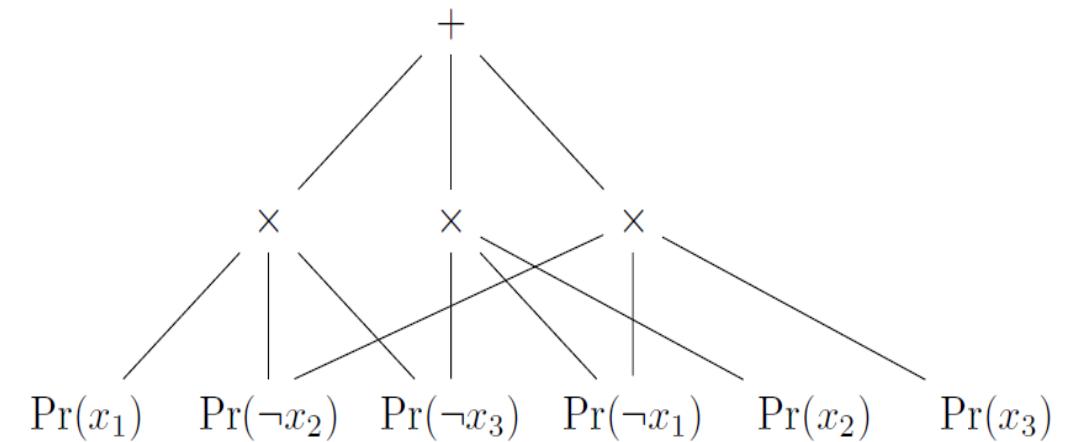
Representation: Prob Exactly One

Learning subject to **Exactly–One**

$$L(\alpha, p) = -\log(p)$$

$\neg Light \wedge Ball$
 $\neg Ball \wedge Light$
Etc.

α

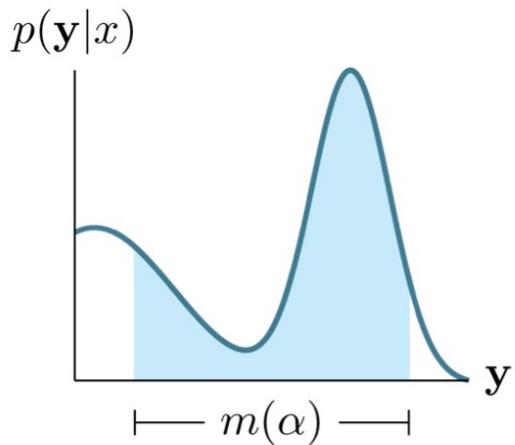


Semantic Loss

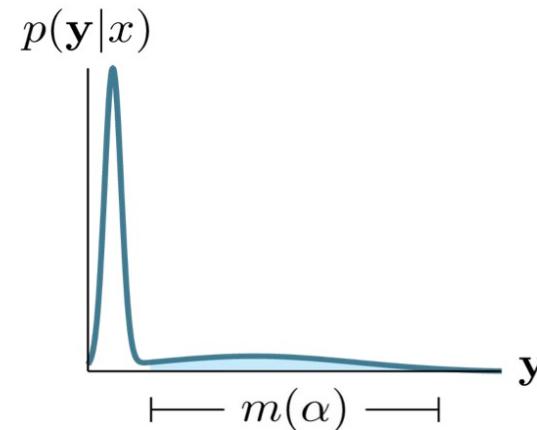
$$L^s(\alpha, p) \propto -\log \sum_{x \models \alpha} \prod_{i: x \models X_i} p_i \prod_{i: x \models \neg X_i} (1 - p_i)$$

- Go over every model that satisfies
(The circuit does it for us)
- Count the probability of every model

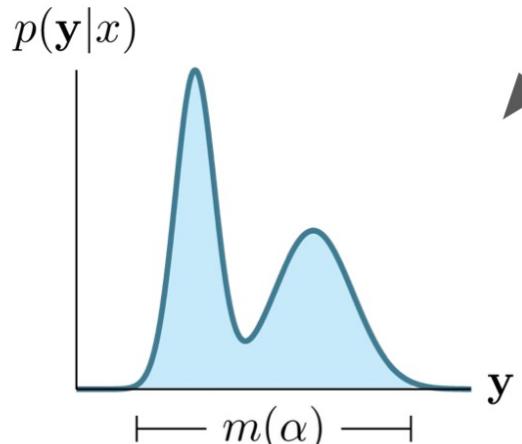
Semantic Loss: Effects



a) A network uncertain over both valid & invalid predictions



b) A network allocating most of its mass to an invalid prediction.

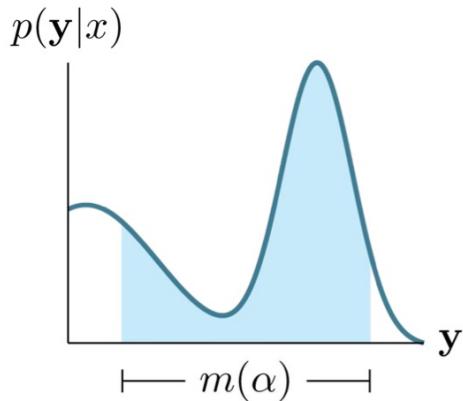


c) A network allocating most of its mass to models of the formula

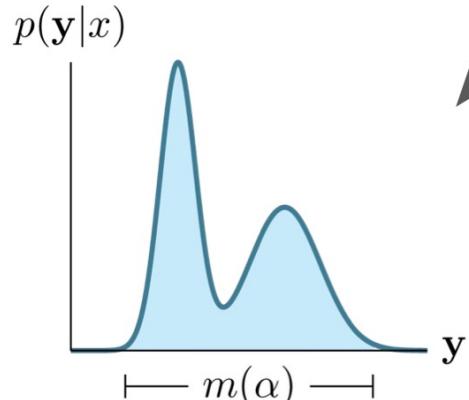
**Semantic
Loss**

$$P(\alpha|x) \uparrow: -\log P(\alpha|x) \downarrow$$

Semantic Loss + Regularization



a) A network uncertain over both valid & invalid predictions

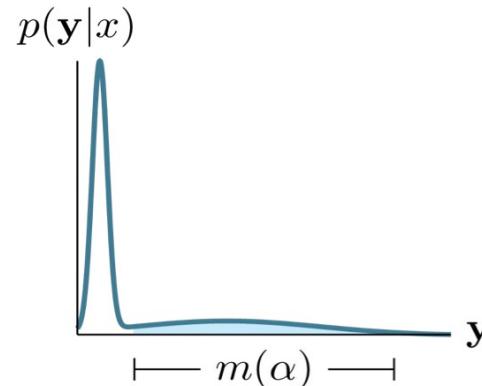


Semantic Loss

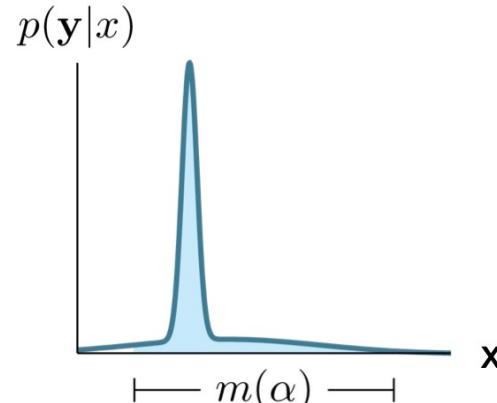
Neuro-Symbolic Entropy Regularization

$$-\mathbb{E}_{P(Y|x,\alpha)}[\log P(Y|x, \alpha)]$$

c) A network allocating most of its mass to models of the formula



b) A network allocating most of its mass to an invalid prediction.



d) A network allocating most of mass to one model of formula

Pylon

Library that extends PyTorch to allow injection of declarative knowledge

- **Easy to Express Knowledge:** users write **arbitrary constraints** on the output
- **Integrates with PyTorch:** **minimal change** to existing code
- **Efficient Training:** compiles into loss that can be **efficiently optimized**
 - Exact semantic loss
 - Monte-carlo estimate of loss
 - T-norm approximation
 - *your solver?*

<http://pylon-lib.github.io>

Pylon

PyTorch Code

```
for i in range(train_iters):
    ...
    py = model(x)
    ...
    loss = CrossEntropy(py,...)
```

1

Specify knowledge as a predicate

```
def check(y):
    ...
    return isValid
```

<http://pylon-lib.github.io>

Pylon

PyTorch Code

```
for i in range(train_iters):
    ...
    py = model(x)
    ...
    loss = CrossEntropy(py, ...)
    loss += constraint_loss(check)(py)
```

1

Specify knowledge as a predicate

```
def check(y):
```

...

```
return isValid
```

2

Add as loss to training

```
loss += constraint_loss(check)
```

<http://pylon-lib.github.io>

Pylon

PyTorch Code

```
for i in range(train_iters):
    ...
    py = model(x)
    ...
    loss = CrossEntropy(py, ...)
    loss += constraint_loss(check)(py)
```

1

Specify knowledge as a predicate

```
def check(y):
```

...

```
return isValid
```

2

Add as loss to training

```
loss += constraint_loss(check)
```

3

pylon derives the gradients
(solves a combinatorial problem)

<http://pylon-lib.github.io>

An Exercise: XOR

PyTorch Code

```
for i in range(train_iters):
    ...
    py = model(x)
    ...
    loss = CrossEntropy(py, ...)
```

1

Specify knowledge as a predicate

```
def check(y):
    ...
    return isValid
```

```
from pylon.constraint import constraint
from pylon.brute_force_solver import SatisfactionBruteForceSolver

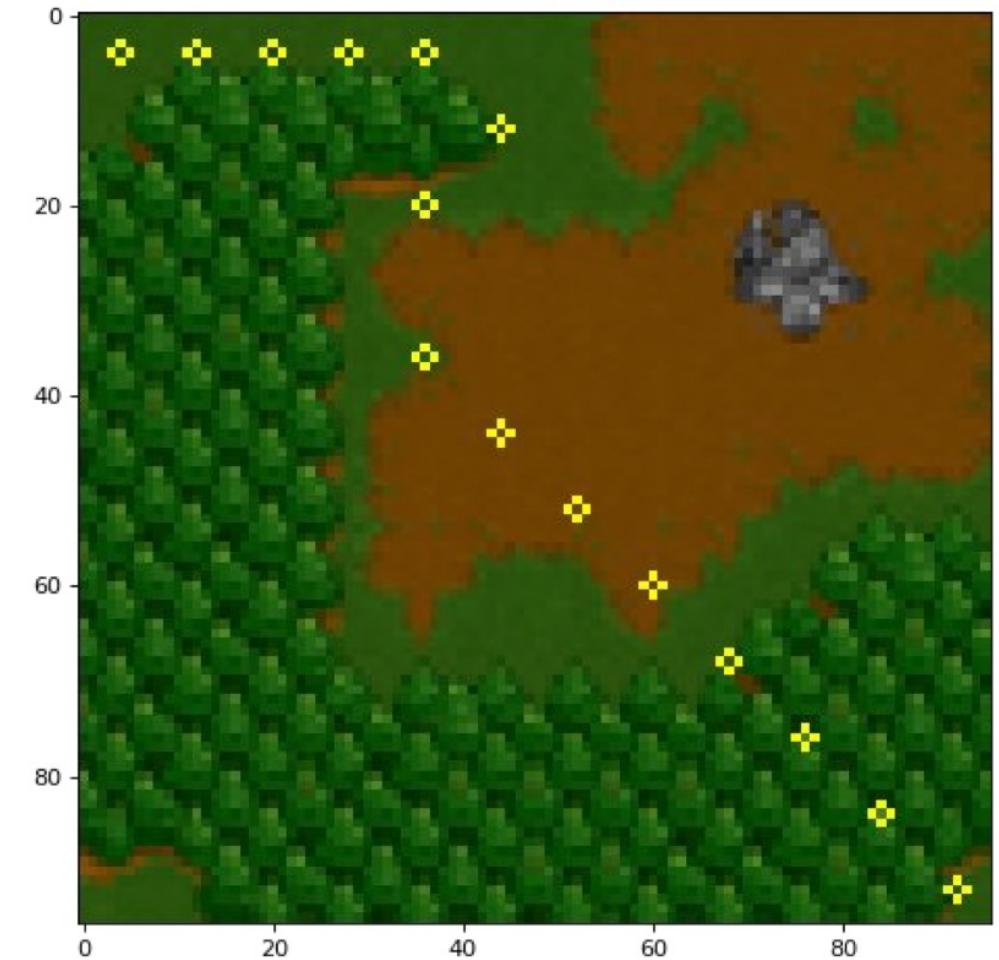
# Our constraint function accepts a decoding tensor of
# shape (batch_size, ...) and is expected to return
# a tensor of shape (batch_size, )
def xor(y):
    return y[:, 0] != y[:, 1] and y[:, 1] != y[:, 2]

xor_cons = constraint(xor, SatisfactionBruteForceSolver())
```

How to extend to exactly one?

<http://pylon-lib.github.io>

An Exercise: Path Constraint



<http://pylon-lib.github.io>

Reasoning with Programming

Task: Classify pairs of MNIST digits with their sum

3	5	8
0	4	4
9	2	11

Reasoning with Programming

Task: Classify pairs of MNIST digits with their sum

3	5	8
0	4	4
9	2	11

Encode single-digit addition as logic

```
addition(X,Y,Z) :- digit(X,N1), digit(Y,N2), Z is N1+N2.
```

Reasoning with Programming

Task: Classify pairs of MNIST digits with their sum

3	5	8
0	4	4
9	2	11

Put neural perception backs

```
nn(mnist_net, [X], Y, [0 ... 9] ) :: digit(X,Y).  
  
addition(X,Y,Z) :- digit(X,N1), digit(Y,N2), Z is N1+N2.
```

Reasoning with Programming

Task: Classify pairs of MNIST digits with their sum

3	5	8
0	4	4
9	2	11

Benefit: separate logic and knowledge

Put neural perception backs

```
nn(mnist_net, [X], Y, [0 ... 9] ) :: digit(X,Y).  
  
addition(X,Y,Z) :- digit(X,N1), digit(Y,N2), Z is N1+N2.
```

An Exercise: Multi-digit Addition

Learn to classify the sum of pairs of MNIST digits

Individual digits are not labeled!

E.g. (**3** , **5** , 8)

Could be done by a CNN: classify the concatenation of both images into 19 classes

However: **3 5 0 4** 1 + **9 2 1** = ?

An Exercise: Multi-digit Addition

Represent multi-digit number

[H | T]

A recursive program

An Exercise: Multi-digit Addition

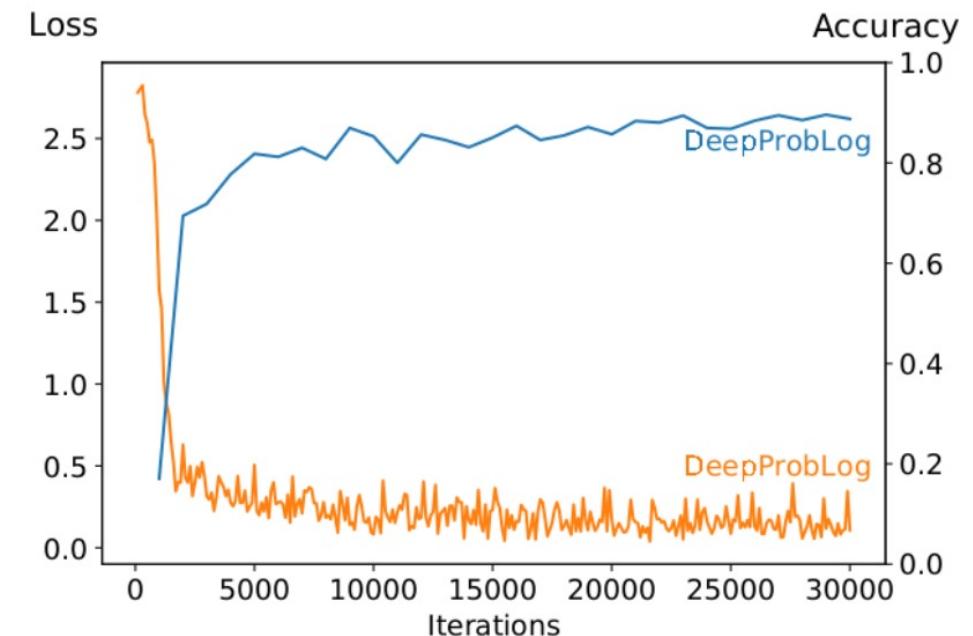
Represent multi-digit number

[H | T]

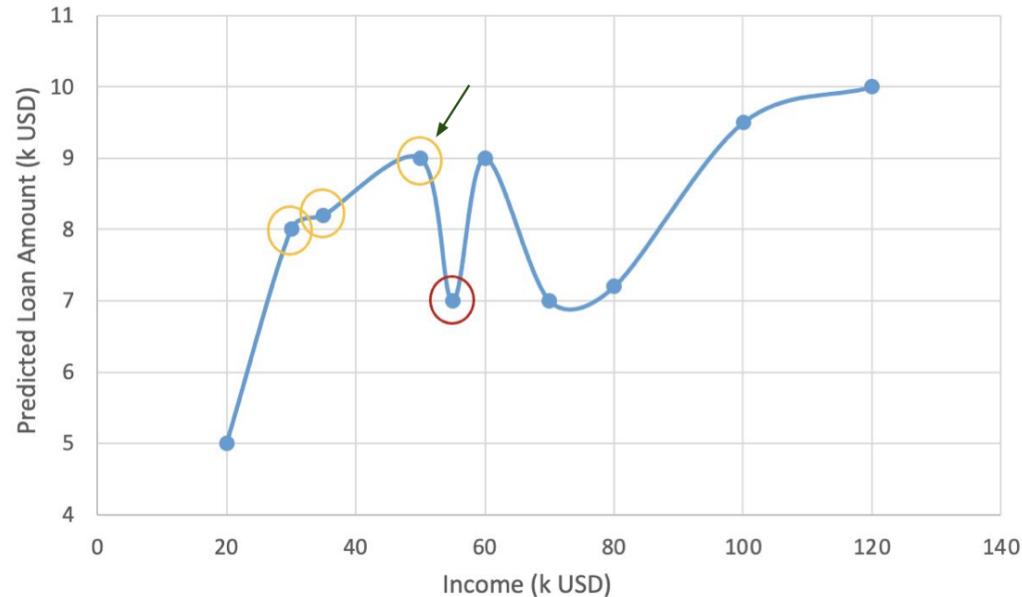
A recursive program

```
number( [ ] , Result , Result ) .  
number( [H | T ] , Acc , Result ) :-  
    digit(H, Nr) , Acc2 is Nr +10*Acc ,  
    number( T , Acc2 , Result ) .  
number( X,Y ) :- number( X, 0 ,Y ) .
```

```
multiaddition(X, Y, Z ) :-  
    number( X, X2 ) ,  
    number( Y, Y2 ) ,  
    Z is X2+Y2 .
```



Reasoning with Counter-Examples



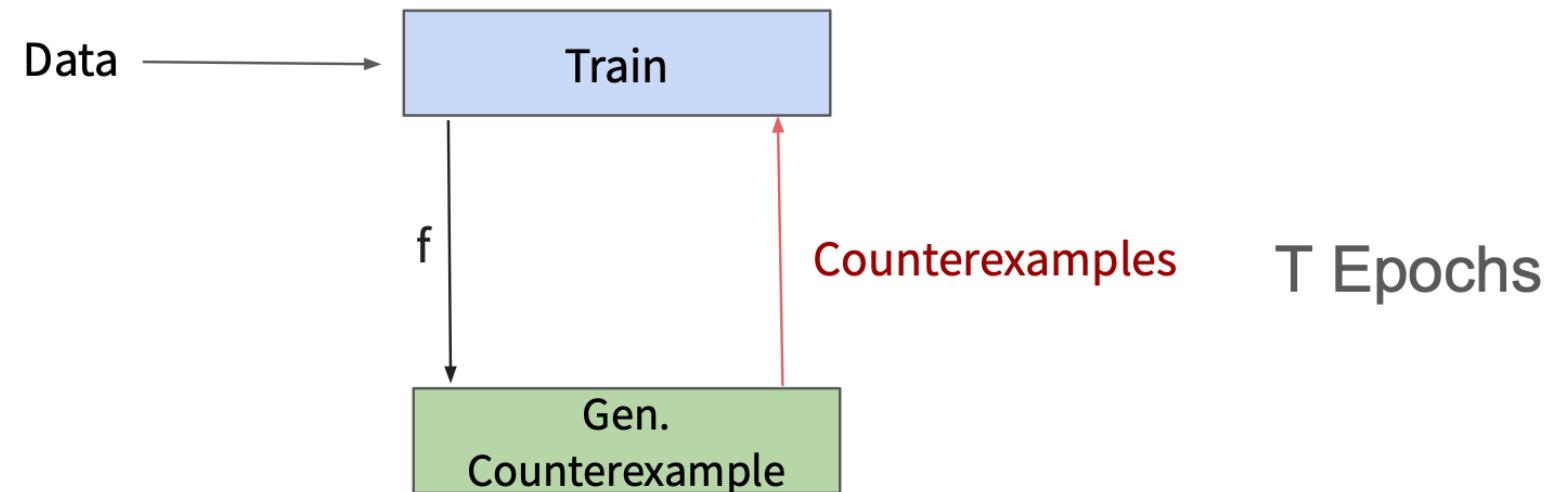
$$\exists x, y \ x \leq y \implies f(x) > f(y)$$

Computed using SMT(LRA)
logical reasoning solver

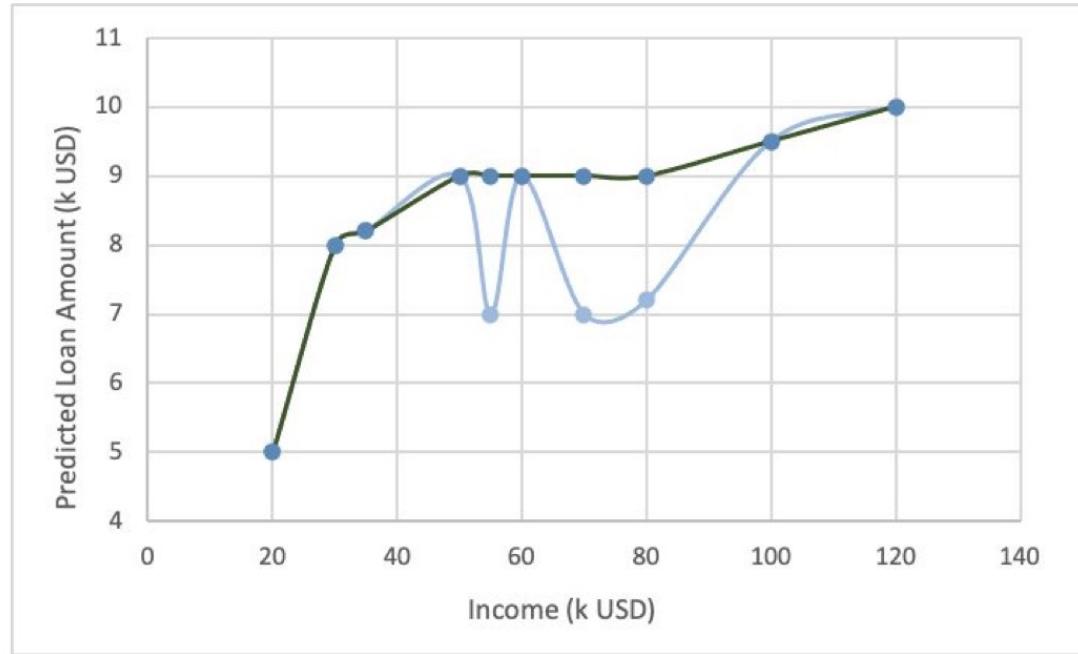
Maximal counterexamples
(largest violation) using OMT

Reasoning with Counter-Examples

How to use monotonicity to improve model quality?
“Monotonicity as inductive bias”



Reasoning with Counter-Examples



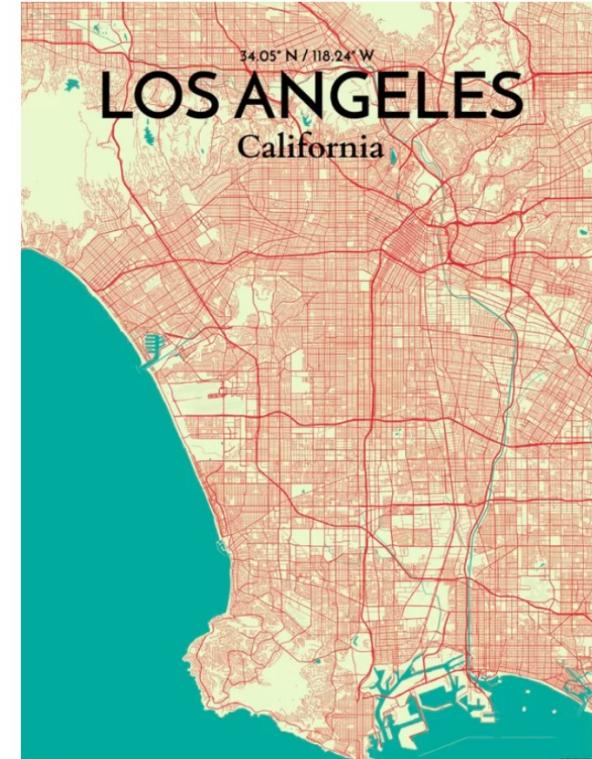
Monotonic Envelope:

- Replace each prediction by its maximal counterexample
- Envelope construction is online (during prediction)
- Guarantees monotonic predictions for any ReLU neural net
- Works for high-dimensional input
- Works for multiple monotonic features

Probabilistic Reasoning

Why Probabilistic Reasoning

q1: *What is the probability that today is a Monday and there is a traffic jam on Westwood Blvd.?*

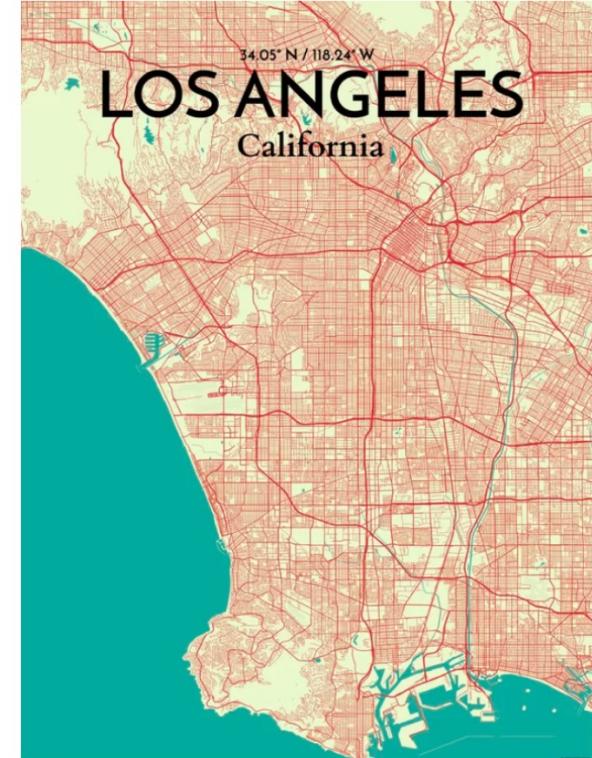


© fineartamerica.com

Why Probabilistic Reasoning

q₁: *What is the probability that today is a Monday and there is a traffic jam on Westwood Blvd.?*

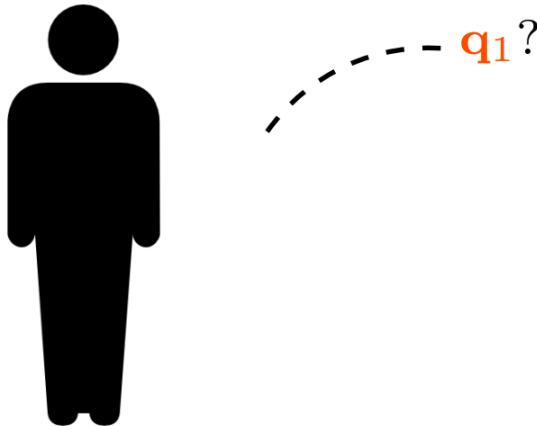
q₂: *Which day is most likely to have a traffic jam on my route to campus?*



© fineartamerica.com

Alternatives

“What is the most likely street to have a traffic jam at 12.00?”

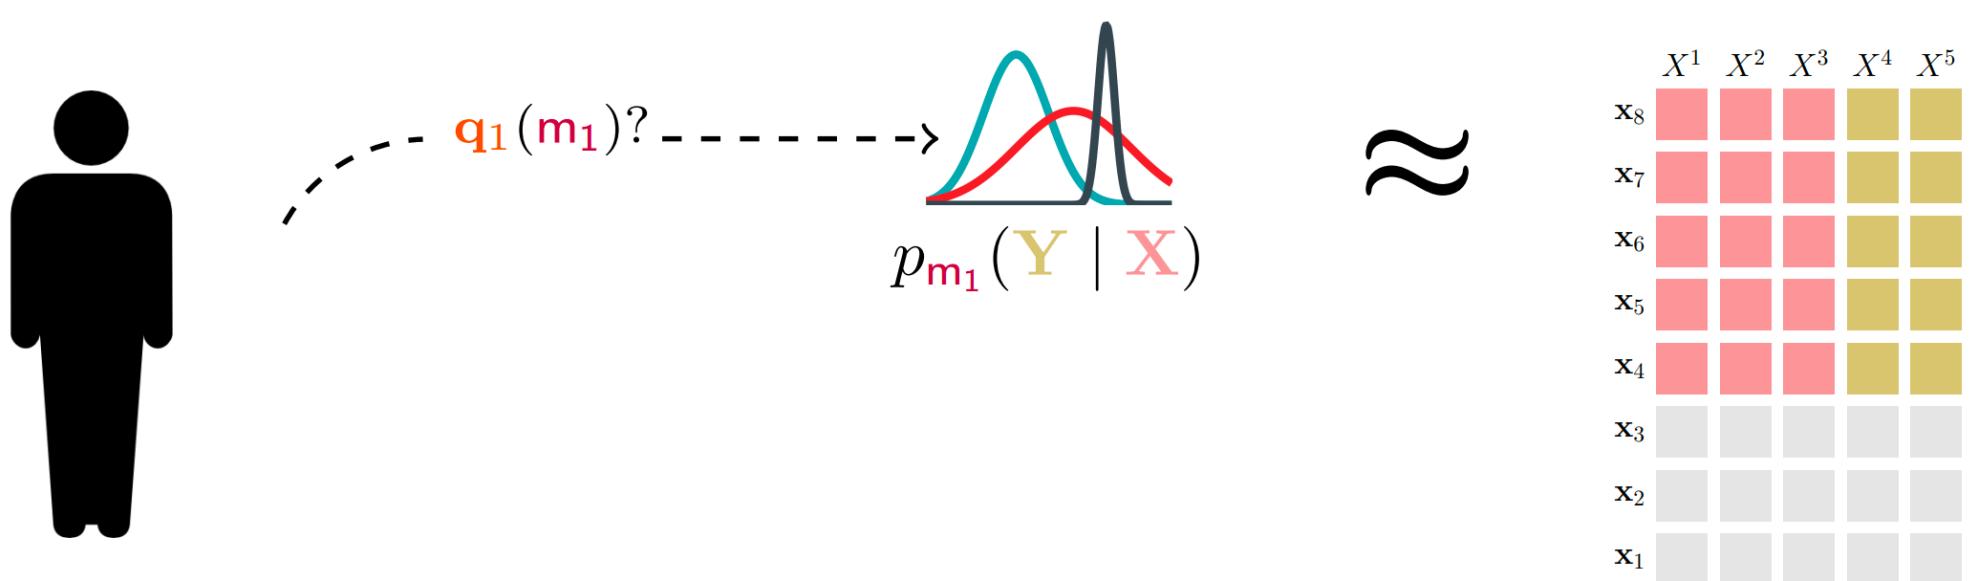


answering queries...

	X^1	X^2	X^3	X^4	X^5
x_8	■	■	■	■	■
x_7	■	■	■	■	■
x_6	■	■	■	■	■
x_5	■	■	■	■	■
x_4	■	■	■	■	■
x_3	■	■	■	■	■
x_2	■	■	■	■	■
x_1	■	■	■	■	■

Alternatives

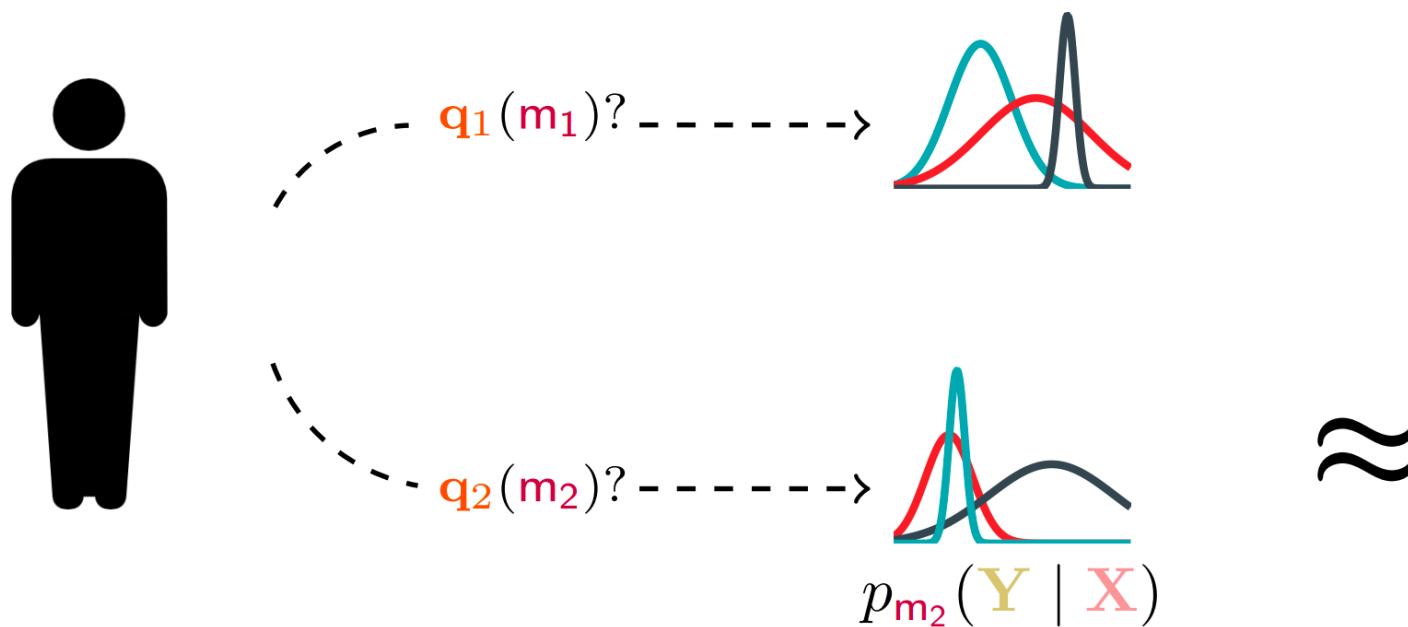
*"What is the most likely **street** to have a traffic jam at **12.00**?"*



...by fitting predictive models!

Alternatives

*“What is the most likely **time** to see a traffic jam at **Sunset Blvd.**? ”*

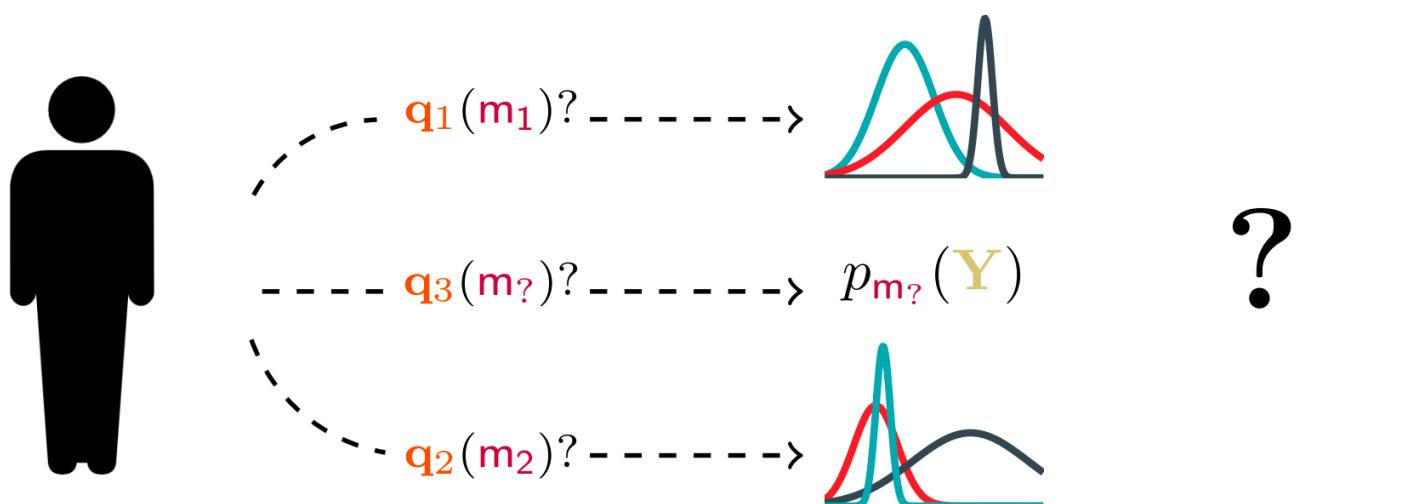


	X^1	X^2	X^3	X^4	X^5
x_8	gray	gray	gray	gray	gray
x_7	gray	gray	gray	gray	gray
x_6	red	red	gold	red	gray
x_5	gray	gray	gray	gray	gray
x_4	red	red	gold	red	gray
x_3	gray	gray	gray	gray	gray
x_2	gray	gray	gray	gray	gray
x_1	red	red	gold	red	gray

~~...by fitting predictive models!~~

Alternatives

"What is the probability of a traffic jam on Westwood Blvd. on Monday?"



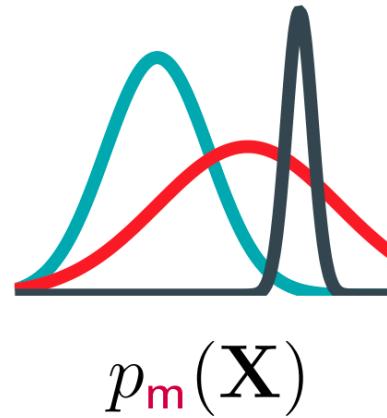
	X^1	X^2	X^3	X^4	X^5
x_8	Gold	Gold	Grey	Grey	Grey
x_7	Gold	Gold	Grey	Grey	Grey
x_6	Gold	Gold	Grey	Grey	Grey
x_5	Gold	Gold	Grey	Grey	Grey
x_4	Gold	Gold	Grey	Grey	Grey
x_3	Gold	Gold	Grey	Grey	Grey
x_2	Gold	Gold	Grey	Grey	Grey
x_1	Gold	Gold	Grey	Grey	Grey

~~...by fitting predictive models!~~

Tractable Probabilistic Reasoning



$q_1(m) ?$
 $q_2(m) ?$
 \dots
 $q_k(m) ?$

A dashed circle surrounds three questions about variables m : $q_1(m)$, $q_2(m)$, and $q_k(m)$. Ellipses between q_2 and q_k indicate additional questions.

\approx

A double approximation symbol (\approx) indicating a relationship or comparison between the generative model and the observed data.

	X^1	X^2	X^3	X^4	X^5
x_8	Red	Red	Red	Green	Green
x_7	Red	Red	Red	Green	Green
x_6	Purple	Yellow	Yellow	Green	Green
x_5	Purple	Blue	Blue	Green	Green
x_4	Purple	Blue	Blue	Green	Green
x_3	Yellow	Yellow	Yellow	Yellow	Yellow
x_2	Brown	Brown	Brown	Brown	Brown
x_1	Brown	Brown	Brown	Brown	Brown

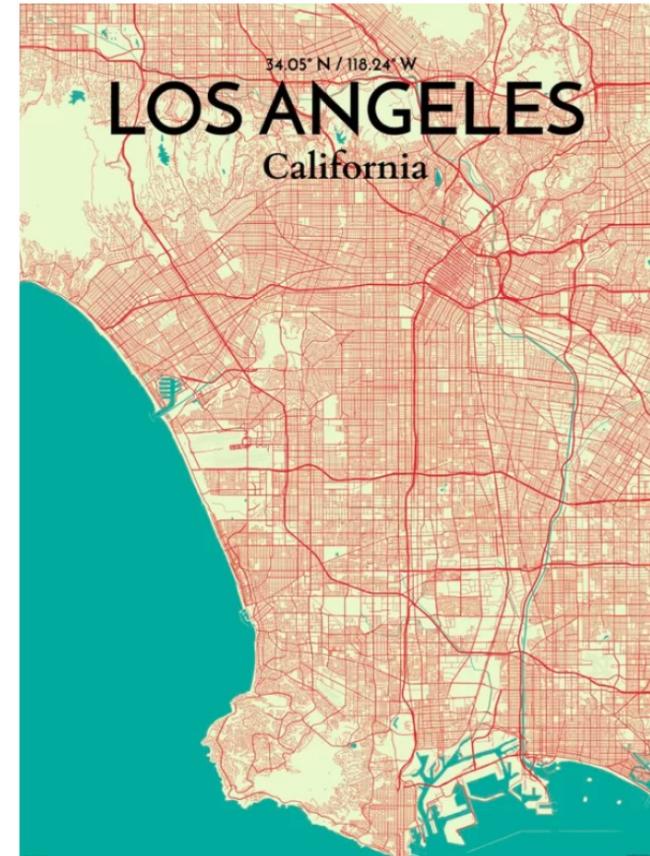
...by fitting generative models!

Probabilistic Reasoning: Complete Evidence

q₃: *What is the probability that today is a Monday at 12.00 and there is a traffic jam only on Westwood Blvd.?*

$$\mathbf{X} = \{\text{Day}, \text{Time}, \text{Jam}_{\text{Westwood}}, \text{Jam}_{\text{Str2}}, \dots, \text{Jam}_{\text{StrN}}\}$$

$$\mathbf{q}_3(\mathbf{m}) = p_{\mathbf{m}}(\mathbf{X} = \{\text{Mon}, 12.00, 1, 0, \dots, 0\})$$



© fineartamerica.com

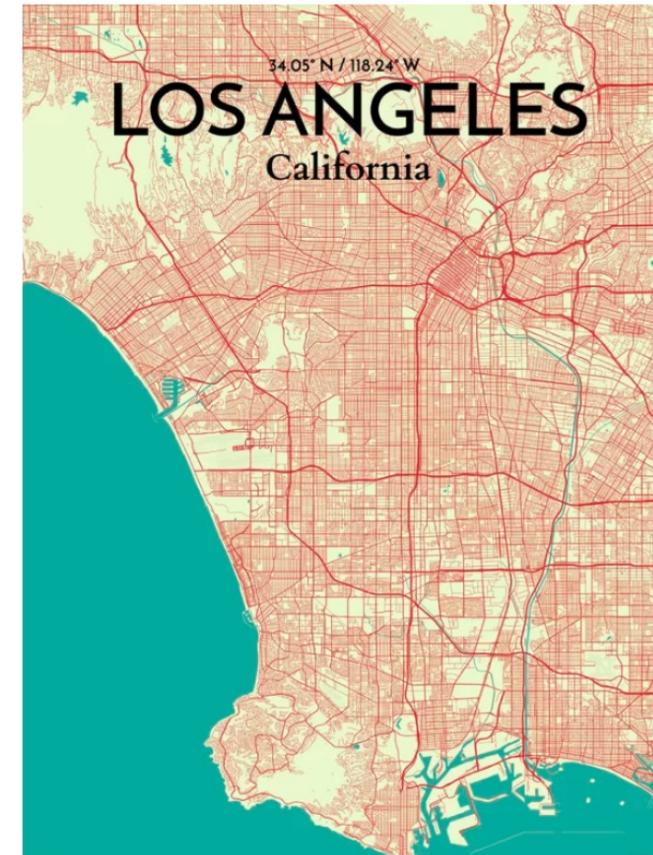
Probabilistic Reasoning: Marginals

q₁: What is the probability that today is a Monday at ~~12:00~~ and there is a traffic jam ~~only~~ on Westwood Blvd.?

$$q_1(\mathbf{m}) = p_{\mathbf{m}}(\text{Day} = \text{Mon}, \text{Jam}_{\text{Westwood}} = 1)$$

$$\text{General: } p_{\mathbf{m}}(\mathbf{e}) = \int p_{\mathbf{m}}(\mathbf{e}, \mathbf{H}) d\mathbf{H}$$

$$\text{where } \mathbf{E} \subset \mathbf{X}, \quad \mathbf{H} = \mathbf{X} \setminus \mathbf{E}$$



© fineartamerica.com

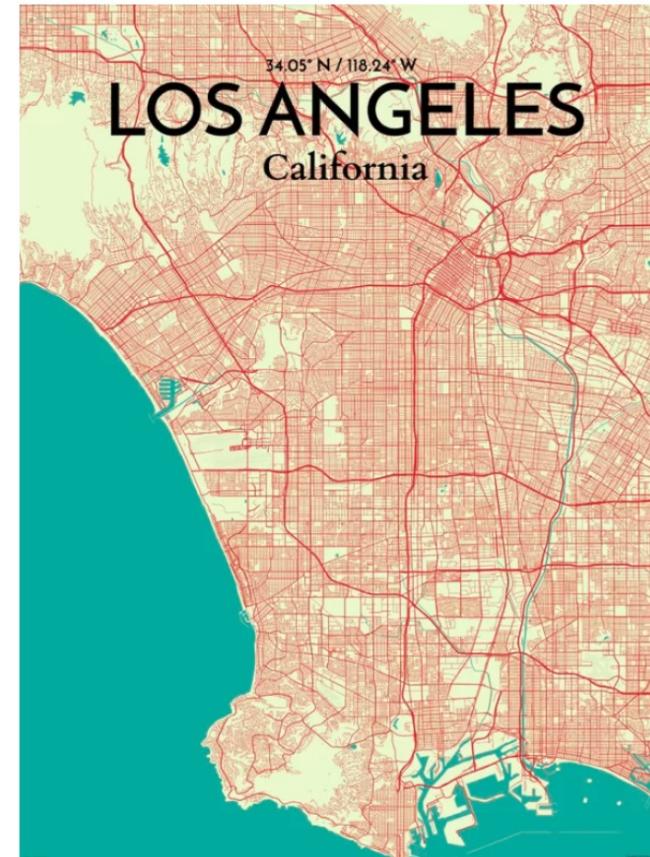
Probabilistic Reasoning: Marginals

q₁: What is the probability that today is a Monday at ~~12:00~~ and there is a traffic jam ~~only~~ on Westwood Blvd.?

$$q_1(\mathbf{m}) = p_{\mathbf{m}}(\text{Day} = \text{Mon}, \text{Jam}_{\text{Westwood}} = 1)$$

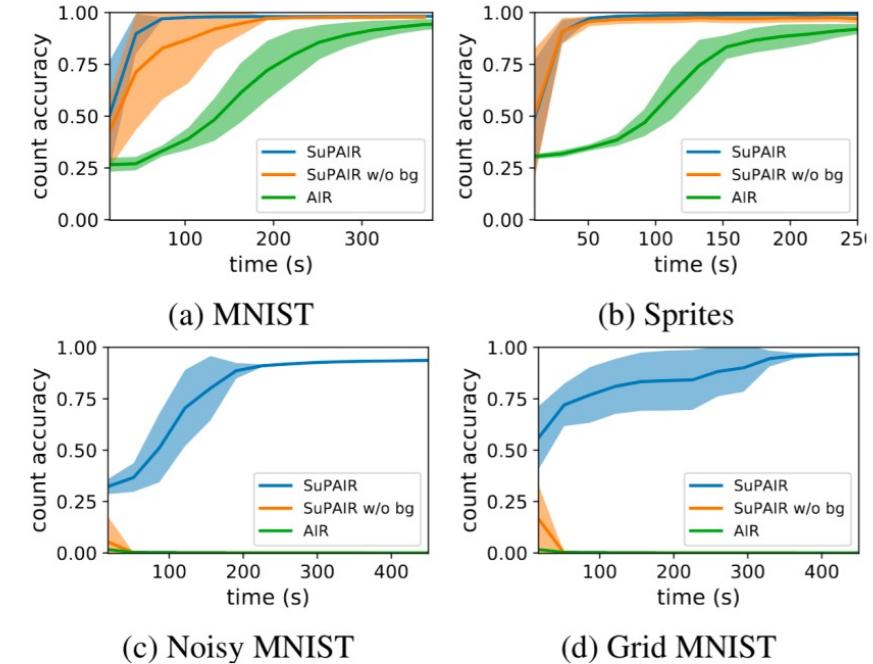
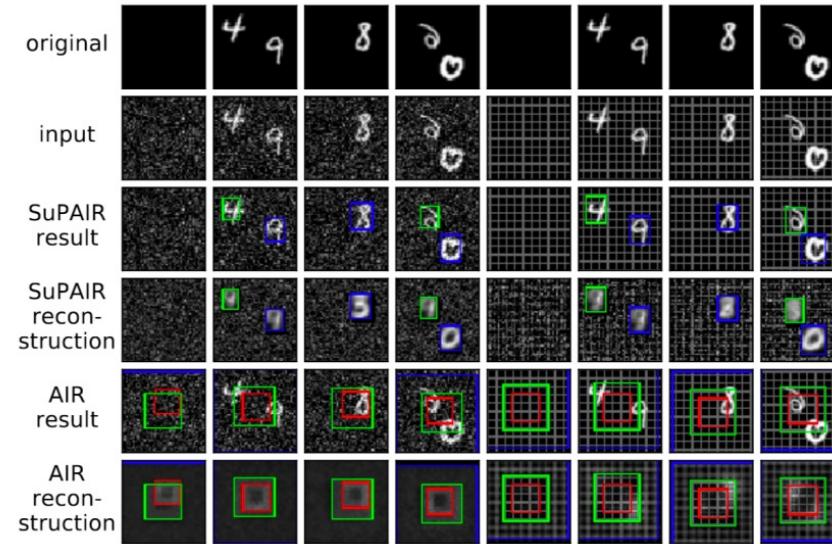
tractable MAR \Rightarrow tractable **conditional queries**
(CON):

$$p_{\mathbf{m}}(\mathbf{q} \mid \mathbf{e}) = \frac{p_{\mathbf{m}}(\mathbf{q}, \mathbf{e})}{p_{\mathbf{m}}(\mathbf{e})}$$



© fineartamerica.com

Probabilistic Reasoning: Marginals



Fast and exact marginalization over unseen or “do not care” parts in the scene

Scene Understanding

Probabilistic Reasoning

A class of queries \mathcal{Q} is tractable on a family of probabilistic models \mathcal{M}

iff for any query $\mathbf{q} \in \mathcal{Q}$ and model $\mathbf{m} \in \mathcal{M}$

exactly computing $\mathbf{q}(\mathbf{m})$ runs in time $O(\text{poly}(|\mathbf{m}|))$.

\Rightarrow often poly will in fact be **linear!**

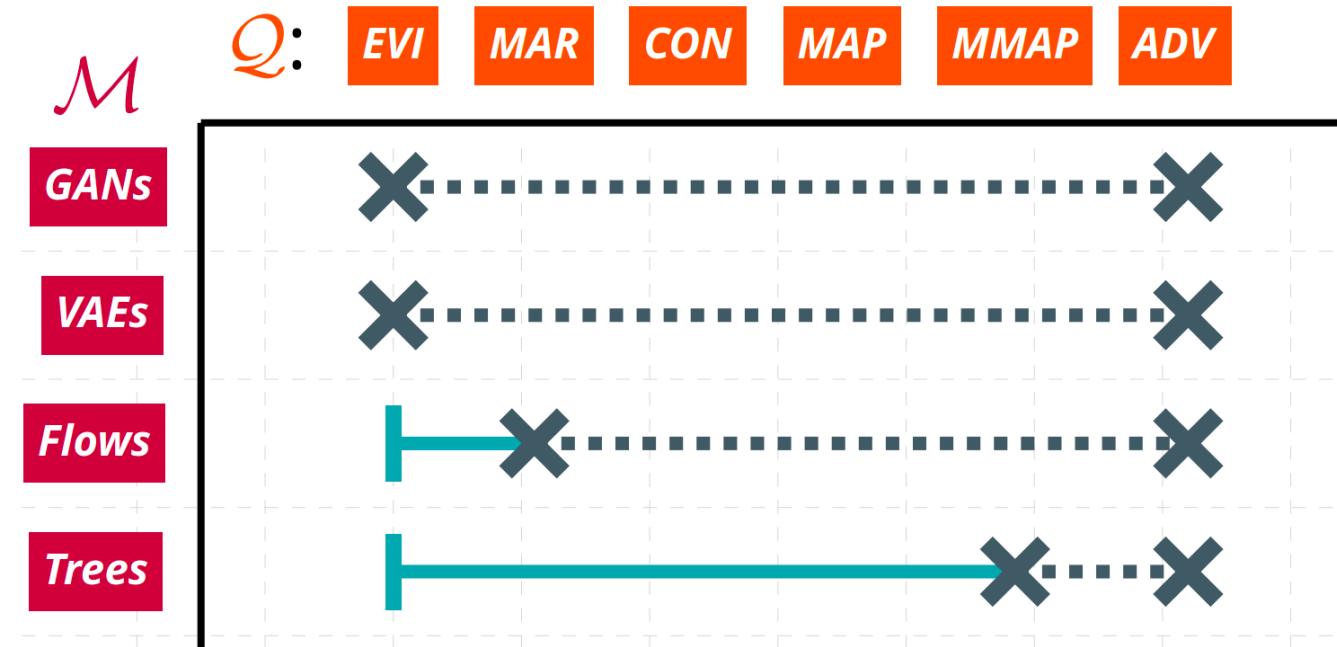
Probabilistic Reasoning

A class of queries \mathcal{Q} is tractable on a family of probabilistic models \mathcal{M} iff for any query $\mathbf{q} \in \mathcal{Q}$ and model $\mathbf{m} \in \mathcal{M}$ exactly computing $\mathbf{q}(\mathbf{m})$ runs in time $O(\text{poly}(|\mathbf{m}|))$.

⇒ often poly will in fact be **linear!**

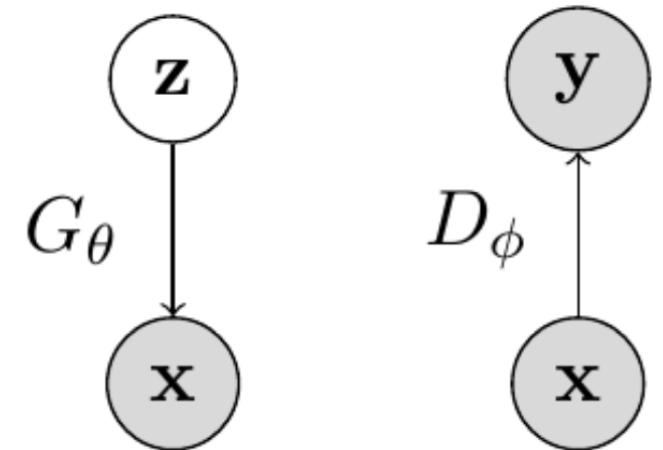
Weighted #SAT is a reasoning

Probabilistic Reasoning: A Spectrum



Probabilistic Reasoning: GANs?

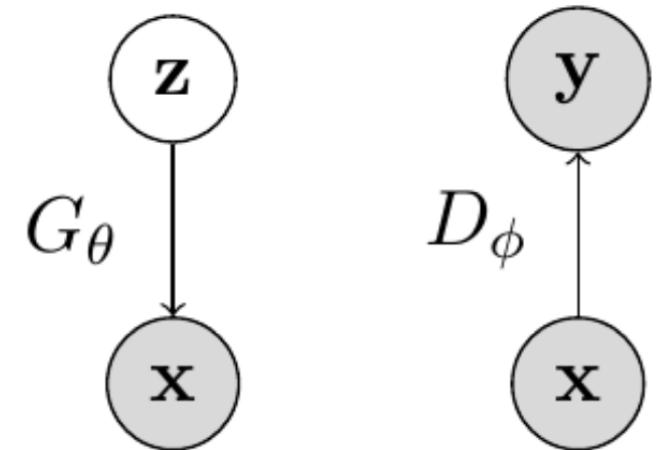
$$\min_{\theta} \max_{\phi} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$



Probabilistic Reasoning: GANs?

$$\min_{\theta} \max_{\phi} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D_{\phi}(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D_{\phi}(G_{\theta}(\mathbf{z})))]$$

- no explicit likelihood!
 - ⇒ *adversarial training instead of MLE*
 - ⇒ *no tractable EVI*
- good sample quality
 - ⇒ *but lots of samples needed for MC*
- unstable training ⇒ *mode collapse*



What Have we Learned from Logical Reasonings

Full factorization: recall AND (decomposability)

A completely disconnected graph. Example: Product of Bernoullis (PoBs)



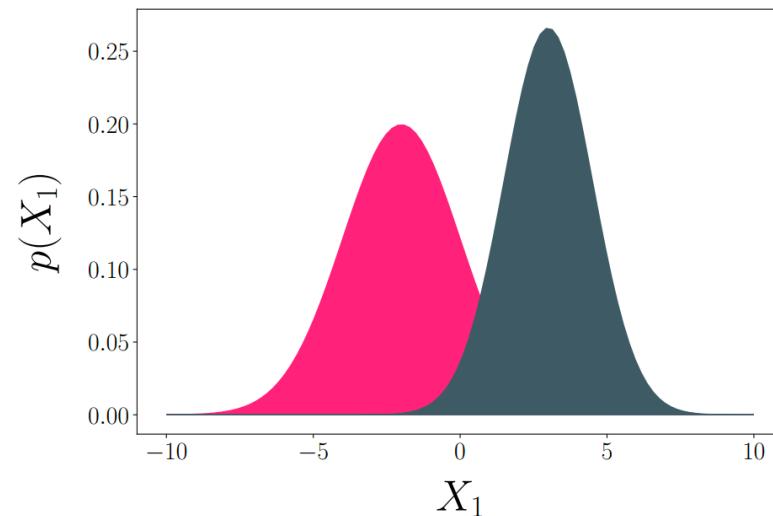
Complete evidence, marginals and MAP, MMAP inference is **linear!**

⇒ *but definitely not expressive...*

What Have we Learned from Logical Reasonings

Mixtures: recall OR (soft determinism)

Mixtures as a convex combination of k (simpler) probabilistic models



$$p(X) = p(Z = 1) \cdot p_1(X|Z = 1) + p(Z = 2) \cdot p_2(X|Z = 2)$$

Mixtures are marginalizing a **categorical latent variable** Z with k values

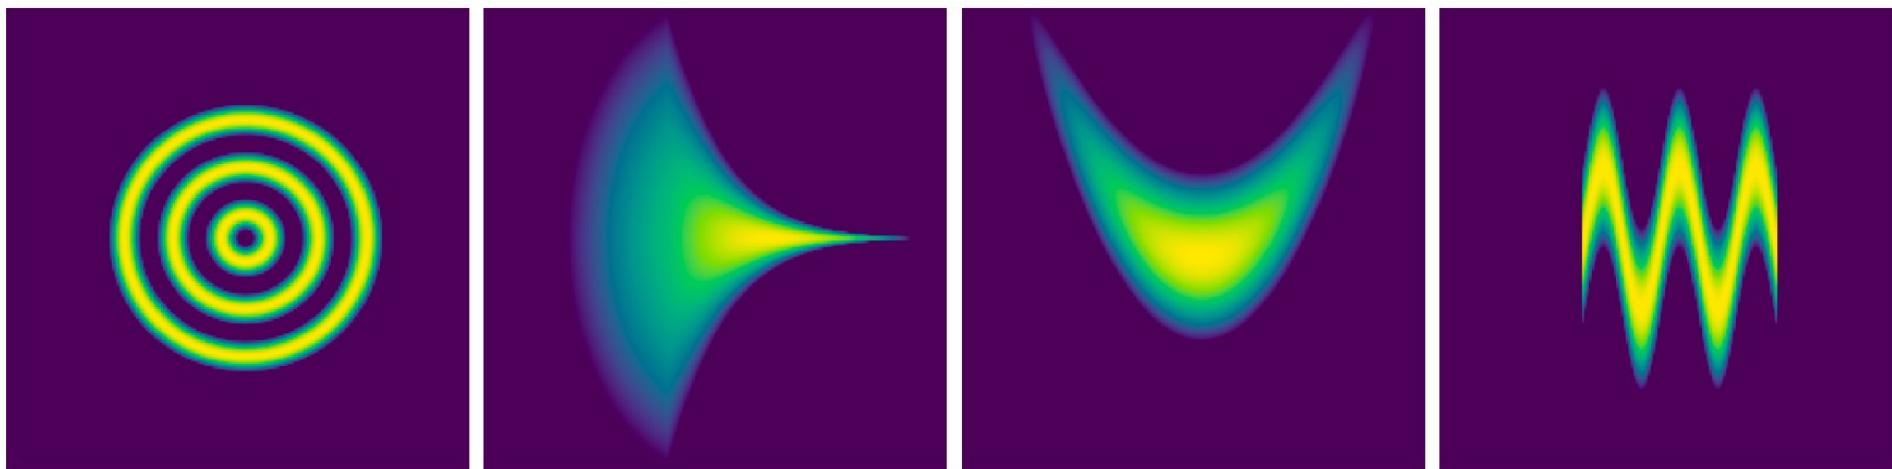
⇒ increased expressiveness

What Have we Learned from Logical Reasonings

Mixtures: recall OR (soft determinism)

Expressiveness: Ability to represent rich and effective classes of functions

⇒ *mixture of Gaussians can approximate any distribution!*

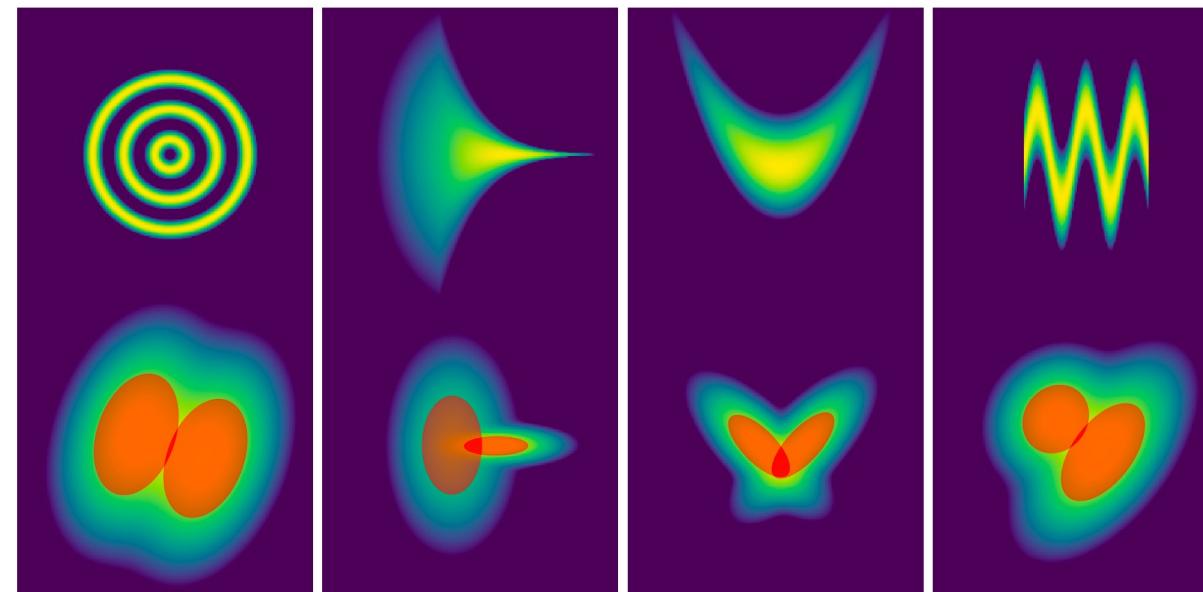


What Have we Learned from Logical Reasonings

Mixtures: recall OR (soft determinism)

Expressiveness: Ability to represent rich and effective classes of functions

⇒ *mixture of Gaussians can approximate any distribution!*

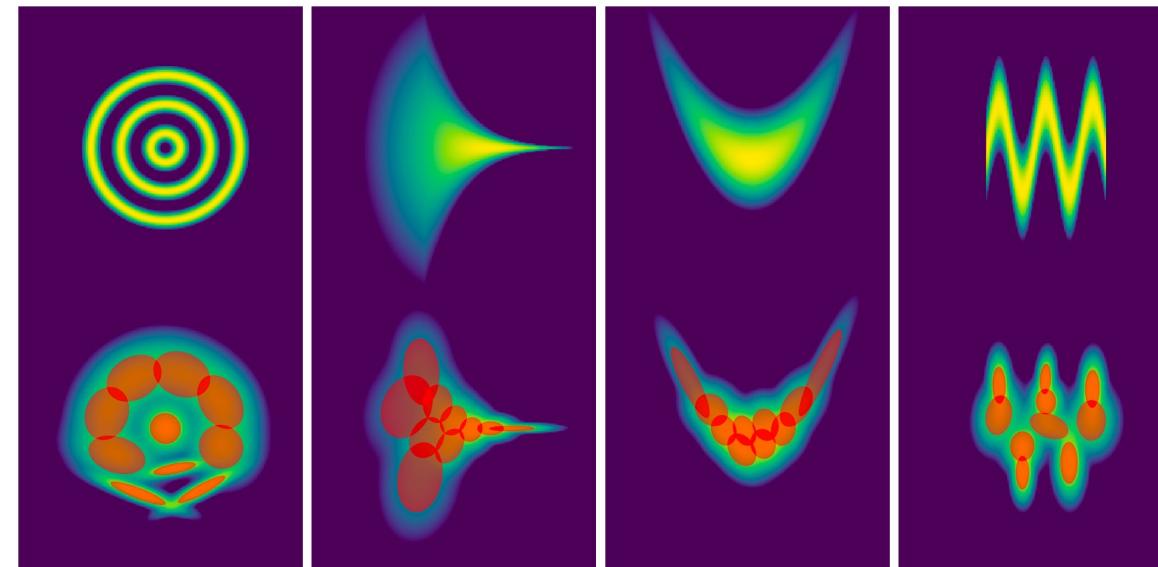


What Have we Learned from Logical Reasonings

Mixtures: recall OR (soft determinism)

Expressiveness: Ability to represent rich and effective classes of functions

⇒ *mixture of Gaussians can approximate any distribution!*

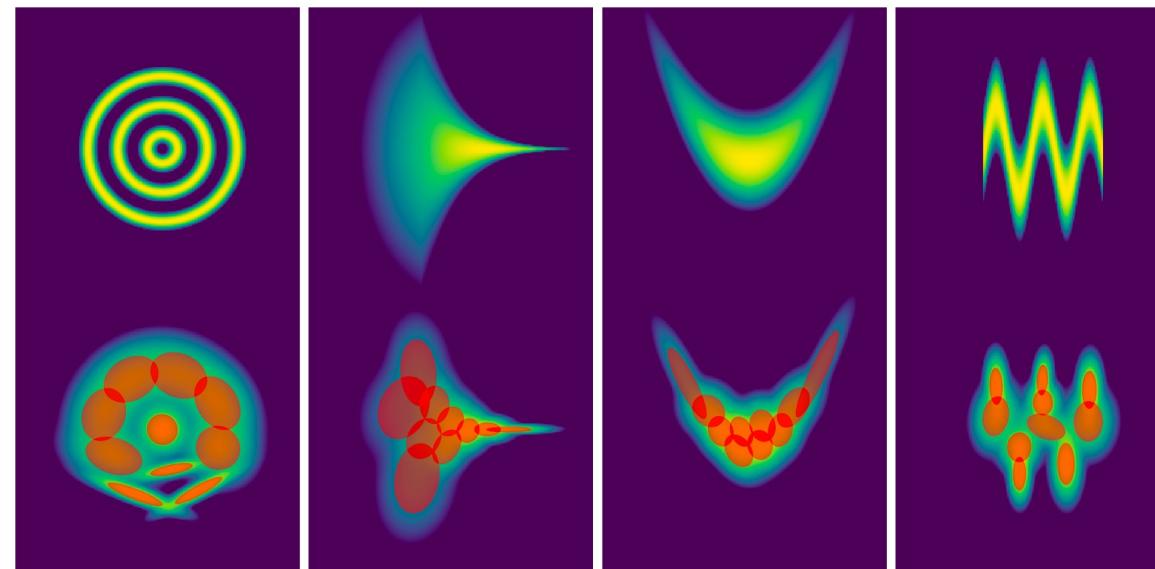


What Have we Learned from Logical Reasonings

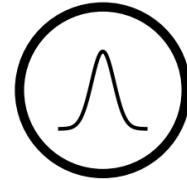
Mixtures: recall OR (soft determinism)

Expressiveness: Ability to represent rich and effective classes of functions

⇒ *mixture of Gaussians can approximate any distribution!*



Tractable Reasoning Model: A Recipe



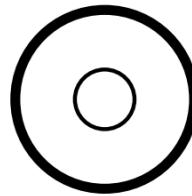
X

Base case: a single node encoding a distribution

⇒ *e.g., Gaussian PDF continuous random variable*

<https://github.com/Juice-jl/>

Tractable Reasoning Model: A Recipe



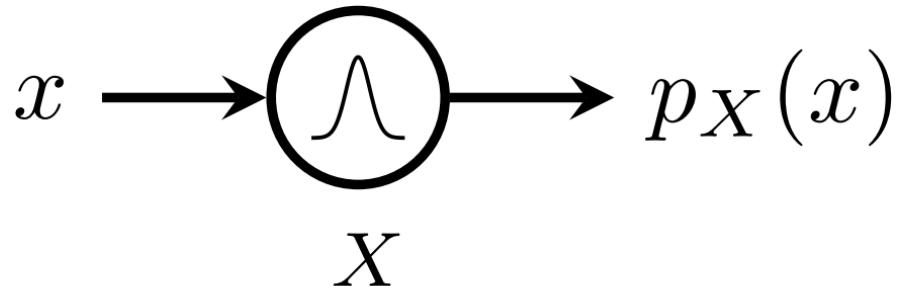
$\neg X$

Base case: a single node encoding a distribution

⇒ e.g., indicators for X or $\neg X$ for Boolean random variable

<https://github.com/Juice-jl/>

Tractable Reasoning Model: A Recipe



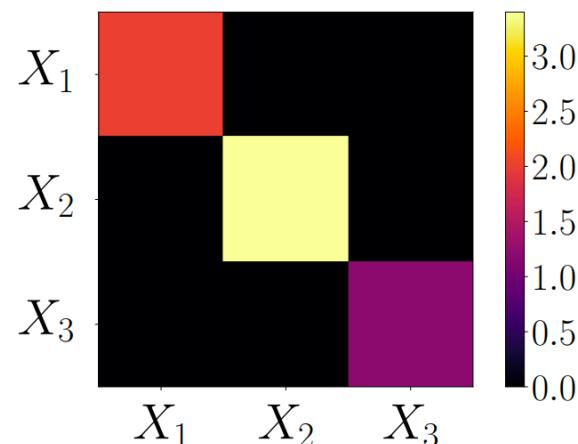
Simple distributions are tractable “black boxes” for:

- EVI: output $p(\mathbf{x})$ (density or mass)
- MAR: output 1 (normalized) or Z (unnormalized)
- MAP: output the mode

Tractable Reasoning Model: A Recipe

Divide and conquer complexity

$$p(X_1, X_2, X_3) = p(X_1) \cdot p(X_2) \cdot p(X_3)$$

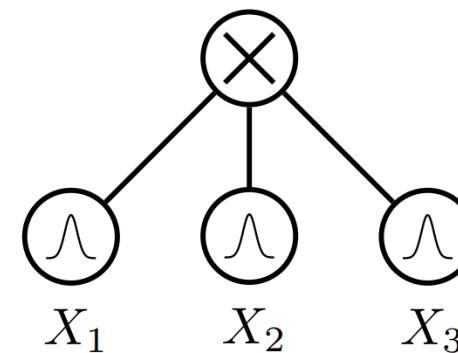
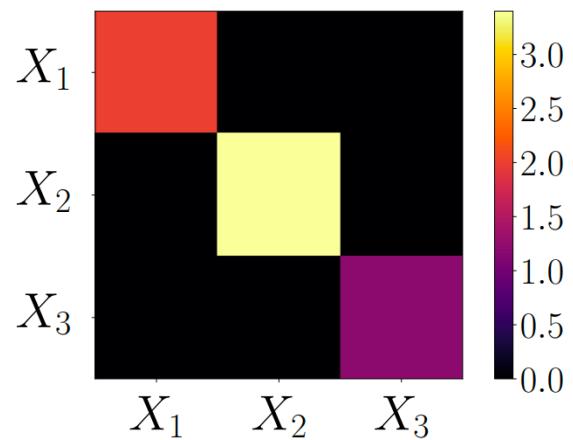


e.g. modeling a multivariate Gaussian with diagonal covariance matrix...

<https://github.com/Juice-jl/>

Tractable Reasoning Model: A Recipe

$$p(X_1, X_2, X_3) = p(X_1) \cdot p(X_2) \cdot p(X_3)$$

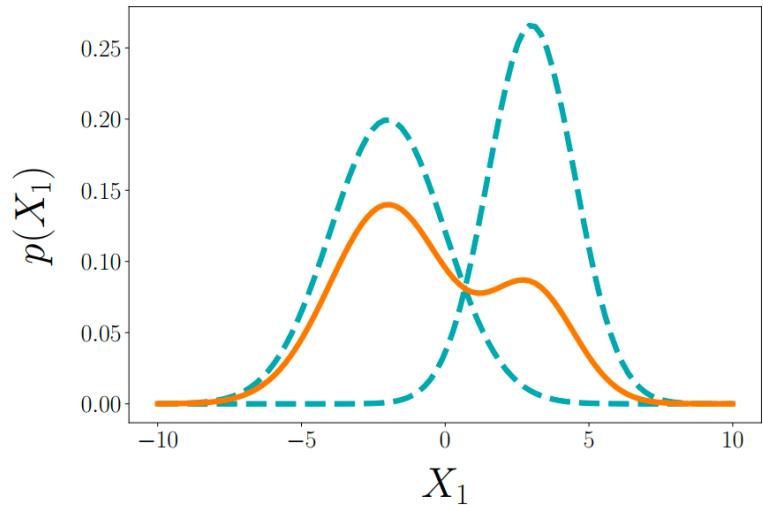


⇒ ...with a product node over some univariate Gaussian distribution

<https://github.com/Juice-jl/>

Tractable Reasoning Model: A Recipe

Enhance expressiveness



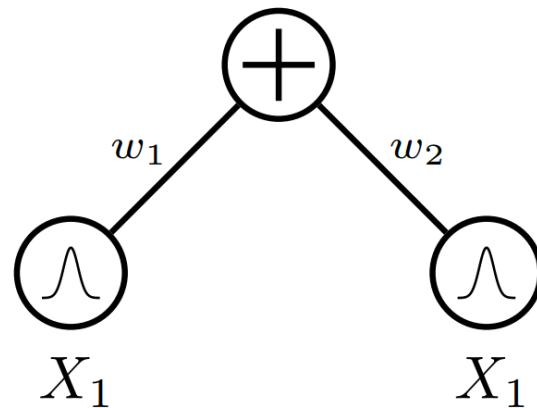
$$p(X) = w_1 \cdot p_1(X) + w_2 \cdot p_2(X)$$

⇒ e.g. modeling a mixture of Gaussians...

<https://github.com/Juice-jl/>

Tractable Reasoning Model: A Recipe

Enhance expressiveness



$$p(x) = 0.2 \cdot p_1(x) + 0.8 \cdot p_2(x)$$

⇒ ...as a weighted sum node over Gaussian input distributions

<https://github.com/Juice-jl/>

Tractable Reasoning Model: A Recipe

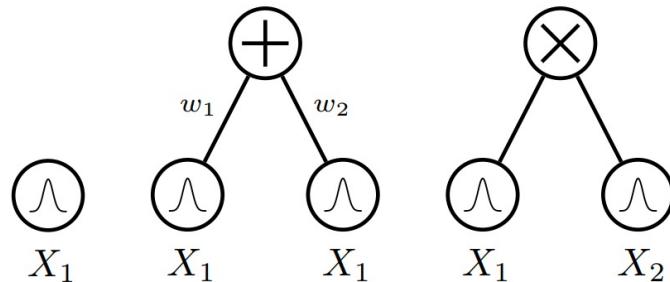
Recursive semantics of probabilistic circuits



<https://github.com/Juice-jl/>

Tractable Reasoning Model: A Recipe

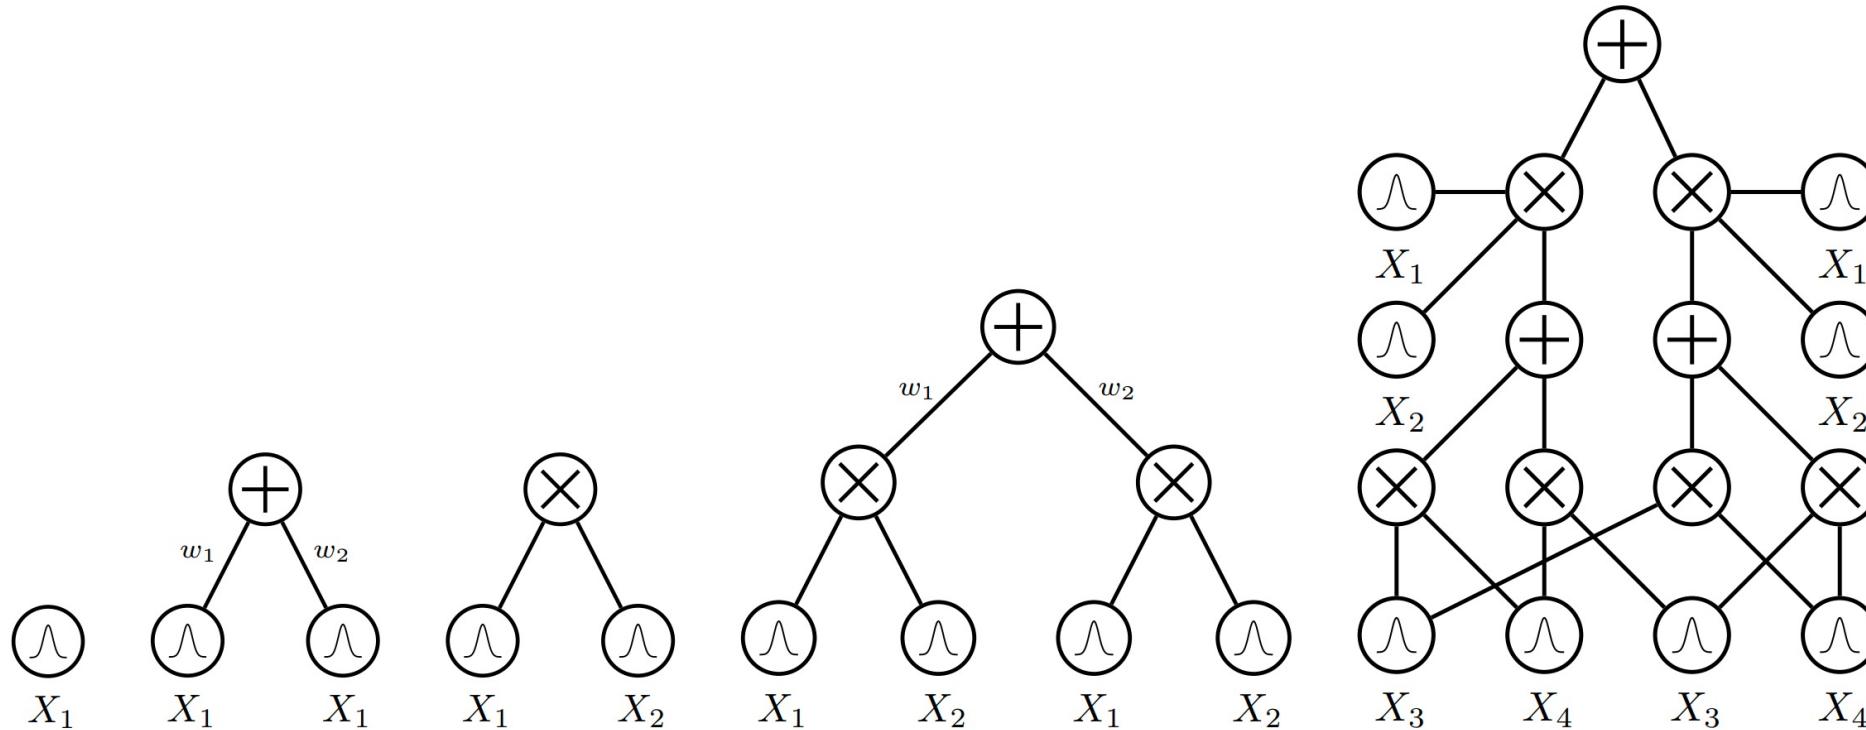
Recursive semantics of probabilistic circuits



<https://github.com/Juice-jl/>

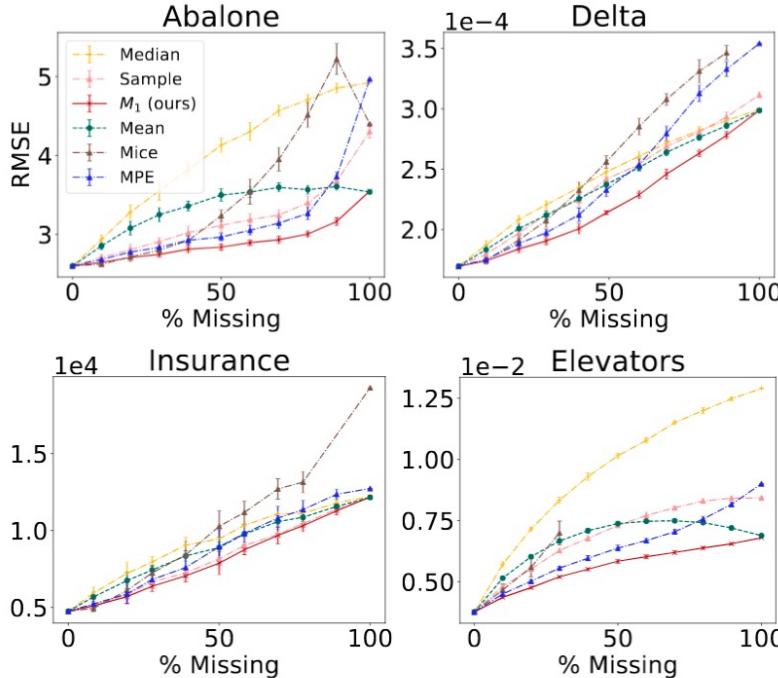
Tractable Reasoning Model: A Recipe

Recursive semantics of probabilistic circuits



<https://github.com/Juice-jl/>

An Example: Expected Predictions



Reasoning about the output of a classifier or regressor f given a distribution p over the input features

$$\mathbb{E}_{\mathbf{x}^m \sim p_\theta(\mathbf{x}^m | \mathbf{x}^o)} [f_\phi^k(\mathbf{x}^m, \mathbf{x}^o)]$$

→ PR if f is a logical formula
→ missing values at test time
→ exploratory classifier analysis

<https://github.com/Juice-jl/>

An Example: Expected Predictions



```
using ProbabilisticCircuits  
pc = load_prob_circuit(zoo_psdd_file("insurance.psdd"));  
rc = load_logistic_circuit(zoo_lc_file("insurance.circuit"), 1);
```

q8: How different is the insurance costs between smokers and non smokers?

```
groups = make_observations([["!smoker"], ["smoker"]])  
exp, _ = Expectation(pc, rc, groups);  
println("Smoker : \$ $(exp[2]);"  
println("Non-Smoker: \$ $(exp[1]);"  
println("Difference: \$ $(exp[2] - exp[1]));"  
Smoker : $ 31355.32630488978  
Non-Smoker: $ 8741.747258310648  
Difference: $ 22613.57904657913
```

<https://github.com/Juice-jl/>

An Example: Expected Predictions



```
using ProbabilisticCircuits  
pc = load_prob_circuit(zoo_psdd_file("insurance.psdd"));  
rc = load_logistic_circuit(zoo_lc_file("insurance.circuit"), 1);
```

q9: Is the predictive model biased by gender?

```
groups = make_observations([["male"], ["female"]])  
exp, _ = Expectation(pc, rc, groups);  
println("Female : \$ $(exp[2])");  
println("Male   : \$ $(exp[1])");  
println("Diff   : \$ $(exp[2] - exp[1])");  
Female : $ 14170.125469335406  
Male   : $ 13196.548926381849  
Diff   : $ 973.5765429535568
```

<https://github.com/Juice-jl/>

Conclusion: the AI Dilemma



- Knowledge is (hidden) everywhere in ML
- A little bit of reasoning goes a long way!

Acknowledgements

Thanks to my collaborators!

Thanks for your attention!

Questions?

